

Introduction à Map-Reduce et Hadoop

Vincent Leroy, **Yves Denneulin,**
Francieli Zanon

Ce que nous allons voir ensemble...

- **Traitement par Batch grande échelle**
 - **paradigme du map/reduce**
 - **Hadoop et HDFS**
- Traitement de données en mémoire
 - Spark
- Cours + TP
 - plus orienté utilisation que performances

Sources

- [Apache Hadoop](#)
- [Yahoo! Developer Network](#)
- [Hortonworks](#)
- [Cloudera](#)
- [Practical Problem Solving with Hadoop and Pig](#)

« Big Data »


- Google, 2008
 - 20 PB/day
 - 180 GB/job (variable)
- Web index
 - 50B pages
 - 15PB
- Criteo : 150 PB

Capacity of a server

- RAM: 256 GB
- Hard drive capacity: 24TB
 - 50PB : ~2100, 150PB : ~6200
- Hard drive throughput: 100MB/s




Data center Google



Pitfalls in parallelism

- Synchronization
 - Mutex,
semaphores ...
- Difficulties
 - Deadlocks
 - Optimization
 - Costly (experts)
 - Not reusable



Fault tolerance

- A server fails every few months
- 1000 servers ...
 - MTBF (mean time between failures) < **1 day**
- A big job may take several days
 - There will be failures, this is **normal**
 - Computations should finish within a reasonable time
 - You cannot start over in case of failures
- Checkpointing, replication
 - Hard to implement correctly at scale

Big Data Platform

- Let everyone write programs for massive datasets
 - Encapsulate parallelism
 - Programming model
 - Deployment
 - Encapsulate fault tolerance
 - Detect and handle failures
- Code once (experts), benefit to all

MAP-REDUCE MODEL

What are Map and Reduce?

- 2 simple functions inspired from functional programming

- **Transformation: map**

$$\text{map}(f, [x_1, \dots, x_n]) = [f(x_1), \dots, f(x_n)]$$

Ex: $\text{map}(*2, [1,2,3]) = [(*2 1), (*2 2), (*2 3)]$
 $= [2,4,6]$

- **Aggregation: reduce**

$$\text{reduce}(f, [x_1, \dots, x_n]) = f(x_1, f(x_2, f(x_3, \dots, f(x_{n-1}, x_n))))$$

Ex: $\text{reduce} (+, [2,4,6]) = (+2 (+4 6))$
 $= 12$

What are Map and Reduce?

- Generic
 - Take a function as a parameter
- Can be **instantiated** and **combined** to solve many different problems
 - `map(toUpperCase, ["hello", "data"])`
= `["HELLO", "DATA"]`
 - `reduce(max, [87, 12, 91])`=91
- The developer provides the functions applied

Data as key/value pairs

- MapReduce does not manipulate atomic pieces of data
 - Everything is a (Key,Value) pair
 - Key and value can be of any type
 - Ex: (Hello, 17)
 - Key = Hello, type text
 - Value = 17 type int
- When initial data is not key/value, interpret it as key/value
 - Input text file becomes [(#line, line_content)...]

Map-Reduce on Key-Value pairs

- Map and Reduce adjusted to Key-Value pairs
 - In map, f is applied **independently** on every key/value pair
$$f(key, value) \rightarrow list(key, value)$$
 - In reduce, f is applied to **all values associated with the same key**
$$f(key, list(value)) \rightarrow list(key, value)$$
 - The types of keys and values taken as input does not have to be the same as the output

Example: Counting frequency of words

- Input : A file of 2 lines
 - 1, "a b c aa b c"
 - 2, "a bb cc a cc b"
- Output
 - a, 3
 - b, 3
 - c, 2
 - aa, 1
 - bb, 1
 - cc, 2

Word frequency: Mapper


- Map processes a portion (line) of text
 - Split words
 - For each word, count one occurrence
 - Key not used in this example (line number)
- ```
map(Int lineNumber, Text line, Output output){
 foreach word in line.split(space) {
 output.write(word, 1)
 }
}
```

# Word frequency: Reducer


- For each key, reduce processes all the corresponding values
  - Add number of occurrences
- ```
reduce(String word, List<Int> occurrences,
Output output){
    int count = 0
    foreach int occ in occurrences {
        count += occ
    }
    output.write(word, count)
}
```

Execution flow

Map



Reduce



How to build a Web index?

- Initial data: (URL, web_page_content)
- Goal: build inverted index

Grenoble

<https://fr.wikipedia.org/wiki/Grenoble>
<http://www.grenoble.fr/>
<http://www.grenoble-tourisme.com/>
<http://wikitravel.org/en/Grenoble>

UNIL

<http://www.unil.ch/>
https://fr.wikipedia.org/wiki/Universit%C3%A9_de_Lausanne
<https://twitter.com/unil>
<http://www.formation-continue-unil-epfl.ch/>

How to build a Web index?

- map(URL pageURL, Text pageContent, Output output){
 foreach word in pageContent.parse() {
 output.write(word, pageURL)
 }
}

How to build a Web index?

- ```
reduce(Text word, List<URL> webPages,
Output output){
 postingList = initPostingList()
 foreach url in webPages {
 postingList.add(url)
 }
 output.write(word, postingList)
}
```


# **APACHE HADOOP: MAPREDUCE FRAMEWORK**

# Objective of Hadoop MapReduce

- Provide a simple and generic programming model: **map** and **reduce**
- Deploy execution automatically
- Provide **fault tolerance**
- **Scale** to thousands of machines
- Performance is important but not the priority
  - What's important is that jobs finish within reasonable time
  - If it's too slow, add servers!  
*Kill It With Iron (**KIWI** principle)*


# Architecture

- From a monolithic architecture to composable layers



# Execution steps

Shuffle & Sort: group by key and transfer to reducer



# Shuffle & Sort

- Barrier in the execution
  - All map tasks must complete before starting reduce
- Partitioner to assign keys to servers executing reduce
  - Ex:  $\text{hash}(\text{key}) \% \text{nbServers}$
  - Deal with load balancing

# Hadoop MapReduce as a developer

- Provide the functions performed by Map and Reduce (Java, C++, Python,...)
  - Application dependent
- Defines the data types (keys / values)
  - If not standard (Text, IntWritable ...)
  - Functions for serialization
- That's all.

# Imports

```
import java.io.IOException ;
import java.util.* ;

import org.apache.hadoop.fs.Path ;
import org.apache.hadoop.io.IntWritable ;
import org.apache.hadoop.io.LongWritable ;
import org.apache.hadoop.io.Text ;
import org.apache.hadoop.mapreduce.Mapper ;
import org.apache.hadoop.mapreduce.Reducer ;
import org.apache.hadoop.mapreduce.JobContext ;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat ;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat ;
import org.apache.hadoop.mapreduce.Job ;
```

**Do not use the old mapred API!**

# Mapper

```
// input key type, input value type, output key type, output
value type
public class WordCountMapper extends Mapper<LongWritable,
Text, Text, IntWritable> {

 @Override
 protected void map(LongWritable key, Text value, Context
context) throws IOException, InterruptedException {
 for (String word : value.toString().split("\\s+")) {
 context.write(new Text(word), new IntWritable(1));
 }
 }
}
```

# Reducer

```
// input key type, input value type, output key type, output
value type
public class WordCountReducer extends Reducer<Text,
IntWritable, Text, LongWritable> {

 @Override
 protected void reduce(Text key, Iterable<IntWritable>
values, Context context) throws IOException,
InterruptedException {
 long sum = 0;
 for (IntWritable value : values) {
 sum += value.get();
 }
 context.write(key, new LongWritable(sum));
 }
}
```

# Main

```
public class WordCountMain {
 public static void main(String [] args) throws Exception {
 Configuration conf = new Configuration();
 String[] otherArgs = new GenericOptionsParser(conf,
args).getRemainingArgs();
 Job job = Job.getInstance(conf, "word count");
 job.setJarByClass(WordCountMain.class);
 job.setMapOutputKeyClass(Text.class);
 job.setMapOutputValueClass(IntWritable.class);
 job.setOutputKeyClass(Text.class);
 job.setOutputValueClass(LongWritable.class);
 job.setMapperClass(WordCountMapper.class);
 job.setReducerClass(WordCountReducer.class);
 job.setInputFormatClass(TextInputFormat.class);
 job.setOutputFormatClass(TextOutputFormat.class);
 FileInputFormat.addInputPath(job, new Path(otherArgs[0]));
 FileOutputFormat.setOutputPath(job, new Path(otherArgs[1]));
 System.exit(job.waitForCompletion(true) ? 0 : 1);
 }
}
```

# Terminology

- MapReduce program = job
- Jobs are submitted to the JobTracker
- A job is divided in several tasks
  - A Map is a task
  - A Reduce is a task
- Tasks are monitored by TaskTrackers
  - A slow task is called a straggler

# Job execution

- $\$ \text{hadoop jar wordcount.jar org.myorg.WordCount}$   
 $\text{inputPath(HDFS) outputPath(HDFS)}$
- Check parameters
  - Is there an output directory ?
  - Does it already exist ?
  - Is there an input directory ?
- Compute splits
- The job (MapReduce code), its configuration and splits are copied with a high replication
- Create an object to follow the progress as the tasks are created by the JobTracker
- For each split, create a Map
- Create default number of reducers

# Tasktracker

- TaskTracker sends a periodic signal to the JobTracker
  - Show that the node still functions
  - Tell whether the TaskTracker is ready to accept a new task
- A TaskTracker is responsible for a node
  - Fixed number of slots for map tasks
  - Fixed number of slots for reduce tasks
  - Tasks can be from different jobs
- Each task runs on its own JVM
  - Prevents a task crash to crash the TaskTracker as well


# Job Progress

- A Map task reports on its progress, i.e. amount of the split processed
- For a reduce task, 3 states
  - copy
  - sort
  - reduce
- Report sent to the TaskTracker
- Every 5 seconds, report forwarded to the JobTracker
- User can see the JobTracker state through Web interface

# Progress

Logged in as: dr.who

## All Applications



Cluster Metrics

| Apps Submitted | Apps Pending | Apps Running | Apps Completed | Containers Running | Memory Used | Memory Total | Memory Reserved | Active Nodes | Decommissioned Nodes | Lost Nodes | Unhealthy Nodes | Rebooted Nodes |
|----------------|--------------|--------------|----------------|--------------------|-------------|--------------|-----------------|--------------|----------------------|------------|-----------------|----------------|
| 1              | 0            | 0            | 1              | 0                  | 0 KB        | 8 GB         | 0 KB            | 1            | 0                    | 0          | 0               | 0              |

Show 20 entries Search:

| ID                             | User   | Name       | Queue   | StartTime               | FinishTime              | State    | FinalStatus | Progress | Tracking UI |
|--------------------------------|--------|------------|---------|-------------------------|-------------------------|----------|-------------|----------|-------------|
| application_1348382082788_0001 | hduser | word count | default | 23-Sep-2012<br>12:06:38 | 23-Sep-2012<br>12:07:20 | FINISHED | SUCCEEDED   |          | History     |

Showing 1 to 1 of 1 entries First Previous 1 Next Last

Cluster

- About
- Nodes
- Applications
  - NEW
  - SUBMITTED
  - ACCEPTED
  - RUNNING
  - FINISHED
  - FAILED
  - KILLED
- Scheduler
- Tools

Switch Theme

[About Apache Hadoop](#)

# End of Job

- Output of each reducer written to a file
- Job tracker notifies the client and writes a report for the job

14/10/28 11:54:25 INFO mapreduce.Job: Job job\_1413131666506\_0070 completed successfully

## Job Counters

Launched map tasks=392

Launched reduce tasks=88

Data-local map tasks=392

[...]

## Map-Reduce Framework

Map input records=622976332

Map output records=622952022

Reduce input groups=54858244

Reduce input records=622952022

Reduce output records=546559709

[...]

## Hadoop job\_200709211549\_0003 on localhost

User: hadoop  
Job Name: streamjob34453.jar  
Job File: /u01/local/hadoop-datasets/hadoop-hadoop-mapreduce-system/job\_200709211549\_0003/job.xml  
Status: Succeeded  
Started at: Fri Sep 21 16:07:10 CEST 2007  
Finished at: Fri Sep 21 16:07:26 CEST 2007  
Finished in: 16sec

| Kind   | % Complete | Num Tasks | Pending | Running | Complete | Killed | Failed/Killed Task Amounts |
|--------|------------|-----------|---------|---------|----------|--------|----------------------------|
| map    | 100.00%    | 3         | 0       | 0       | 3        | 0      | 0 / 0                      |
| reduce | 100.00%    | 1         | 0       | 0       | 1        | 0      | 0 / 0                      |

|                      | Counter               | Map       | Reduce  | Total     |
|----------------------|-----------------------|-----------|---------|-----------|
| Job Counters         | Launched map tasks    | 0         | 0       | 3         |
|                      | Launched reduce tasks | 0         | 0       | 1         |
|                      | Data-local map tasks  | 0         | 0       | 3         |
| Map-Reduce Framework | Map input records     | 77,637    | 0       | 77,637    |
|                      | Map output records    | 103,939   | 0       | 103,939   |
|                      | Map input bytes       | 3,659,910 | 0       | 3,659,910 |
|                      | Map output bytes      | 1,083,757 | 0       | 1,083,757 |
|                      | Reduce input groups   | 0         | 85,095  | 85,095    |
|                      | Reduce input records  | 0         | 103,909 | 103,909   |
|                      | Reduce output records | 0         | 85,095  | 85,095    |

Change priority from NORMAL to: [VERY HIGH](#) [HIGH](#) [LOW](#) [VERY LOW](#)

# Server failure during a job

- Bug in a task
  - task JVM crashes → TaskTracker JVM notified
  - task removed from its slot
- Task become unresponsive
  - timeout after 10 minutes
  - task removed from its slot
- Each task may be re-run up to N times (default 7) in case of crashes

# **HDFS : DISTRIBUTED FILE SYSTEM**

# Random vs Sequential disk access

- Example
    - DB 100M users
    - 100B/user
    - Alter 1% records
  - Random access
    - Seek, read, write: 30mS
    - 1M users → 8h20
  - Sequential access
    - Read ALL Write ALL
    - 2x 10GB @ 100MB/S → 3 minutes
- It is often faster to read all and write all sequentially

# Distributed File System (HDFS)

- Goal
  - Fault tolerance (redundancy)
  - Performance (parallel access)
- Large files
  - Sequential reads
  - Sequential writes
- “in place” data processing
  - Data is stored on the machines that process it
    - Better usage of machines (no dedicated *filer*)
    - Less network bottlenecks (better performance)

# HDFS model

- Data organized in files and directories
  - mimics a standard file system
- Files divided in blocks (default: 64MB) spread on servers
- HDFS reports the data layout to the Map-Reduce framework
  - If possible, process data on the machines where it is already stored

# Fault tolerance

- File blocks replicated (default: 3) to tolerate failures
- Placement according to different parameters
  - Power supply
  - Network equipment
  - Diverse servers to increase the probability of having a “close” copy
- Checksum of data to detect corrupter blocks (also available in modern file systems)

# Master/Worker architecture

- A *master*, the NameNode
  - Manage the space of file names
  - Manages access rights
  - Supervise operations on files, blocks ...
  - Supervise the *health* of the file system (failures, load balance...)
- Many (1000s) workers, the DataNodes
  - Store the data (blocks)
  - Perform read and write operations
  - Perform copies (replication, ordered by the NameNode)

# NameNode

- Stores the metadata of each file and block (*inode*)
  - File name, directory, blocks associated, position of these blocks, number of replicas ...
- Keeps all in main memory (RAM)
  - Limiting factor = number of files
  - 60M objects in 16GB

# DataNode

- Manage and monitor the state of blocks stored on the host file system (often Linux)
- Directly accessed by the clients → data never transit through the NameNode
- Send *heartbeats* to the NameNode to show that the server has not failed
- Report to the NameNode if blocks are corrupted


# Writing a file

- The client sends a query to the NameNode to create a new file
- The NameNode checks
  - Client authorizations
  - File system conflicts (existing file ...)
- NameNode chooses DataNodes to store file and replicas
  - DataNodes “pipelined”
- Blocks are allocated on these DataNodes
- Stream of data sent to the first DataNode of the pipeline
- Each DataNode forwards the data received to the next DataNode in the pipeline

# Reading a file

- Client sends a request to the NameNode to read a file
- NameNode checks the file exists and builds a list of DataNodes containing the first blocks
- For each block, NameNode sends the address of the DataNodes hosting them
  - List ordered wrt. Proximity to the client
- Client connects to the closest DataNode containing the 1<sup>st</sup> block of the file
- Block read ends:
  - Close connection to the DataNode
  - New connection to the DataNode containing the next block
- When all blocks are read:
  - Query the NameNode to retrieve the following blocks

# HDFS Structure



# HDFS commands (directories)

- Create directory dir

*\$ hadoop dfs -mkdir /dir*

- List HDFS content

*\$ hadoop dfs -ls*

- Remove directory dir

*\$ hadoop dfs -rmdir /dir*

# HDFS commands (files)

- Copy local file `toto.txt` to HDFS dir/  
`$ hadoop dfs -put toto.txt dir/toto.txt`
- Copy HDFS file to local disk  
`$ hadoop dfs -get dir/toto.txt ./`
- Read file `/dir/toto.txt`  
`$ hadoop dfs -cat /dir/toto.txt`
- Remove file `/dir/toto.txt`  
`$ hadoop dfs -rm /dir/toto.txt`

# Quelques compléments

**That is costly!**




Figure from <https://www.supinfo.com/articles/single/2807-introduction-to-the-mapreduce-life-cycle>

55 of 61


|                 |
|-----------------|
| 1, "aaa bb ccc" |
| 2, "bb bb d"    |
| 3, "d aaa bb"   |
| 4, "d"          |


```
map(key, value):
 for each word in value:
 output pair(word, 1)
```


|          |
|----------|
| "aaa", 1 |
| "bb", 1  |
| "ccc", 1 |
| "bb", 1  |
| "bb", 1  |
| "d", 1   |
| "d", 1   |
| "aaa", 1 |
| "bb", 1  |
| "d", 1   |

|          |
|----------|
| "aaa", 2 |
| "bb", 4  |
| "ccc", 1 |
| "d", 3   |

```
reduce(key, values):
 result = 0
 for value in values:
 result += value
 output pair(key, result)
```







# Combiner

- Same API as reduce (key, List<value>)
  - Not the same contract!  
For one key, you get SOME values
- Often the same aggregation as reduce
  - E.g. WordCount
- Different when using global properties
  - E.g. Keep words present at least 5 times

# Le TP

- Prise en main de hadoop et map/reduce en programmant en Python
  - Beaucoup de lignes de commande sous Linux, maîtrise nécessaire
- Deux parties
  - écriture et exécution de wordcount
  - analyse de données Flickr
- Exécution sur un cluster Hadoop dédié
  - accessible depuis le réseau Ensimag
- Travail principal : écrire les programmes en charge de la phase map et de la phase reduce
  - 2 programmes à écrire map.py reduce.py
  - entrée mapper -> ligne de fichiers, sortie mapper : clé \t valeur
  - entrée reducer : liste de couples (clé \t valeur) triés par valeur de clé
- Comptes sur le cluster Hadoop -> votre login Ensimag/MS25
- Question à se poser : est-ce que mon code marche si le fichier en entrée fait 10To ?

