

Interactive computations and outsourcing

Jean-Louis Roch
Grenoble University, France

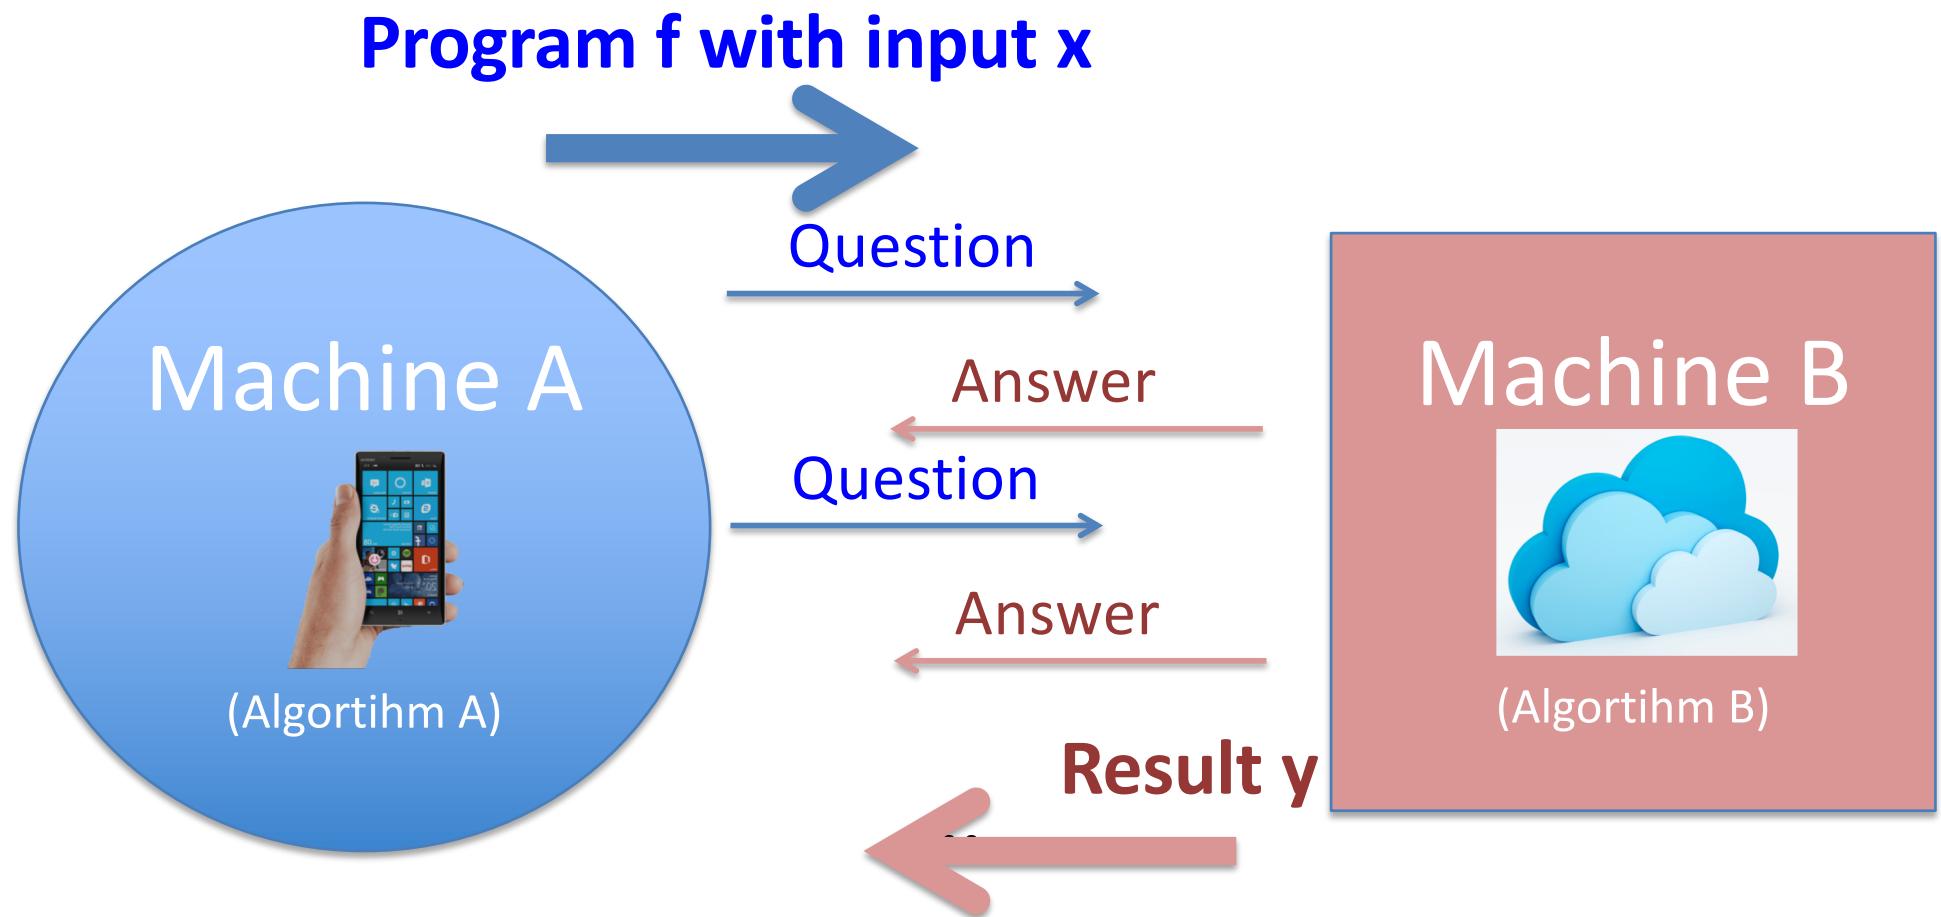
***Grenoble INP -Ensimag, Univ. Grenoble
Alpes, 2021/10/22***

Distributed, heterogeneous (hybrid)



- Various computing abilities and levels of trust

Outsourcing



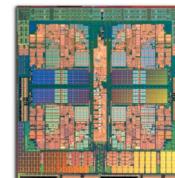
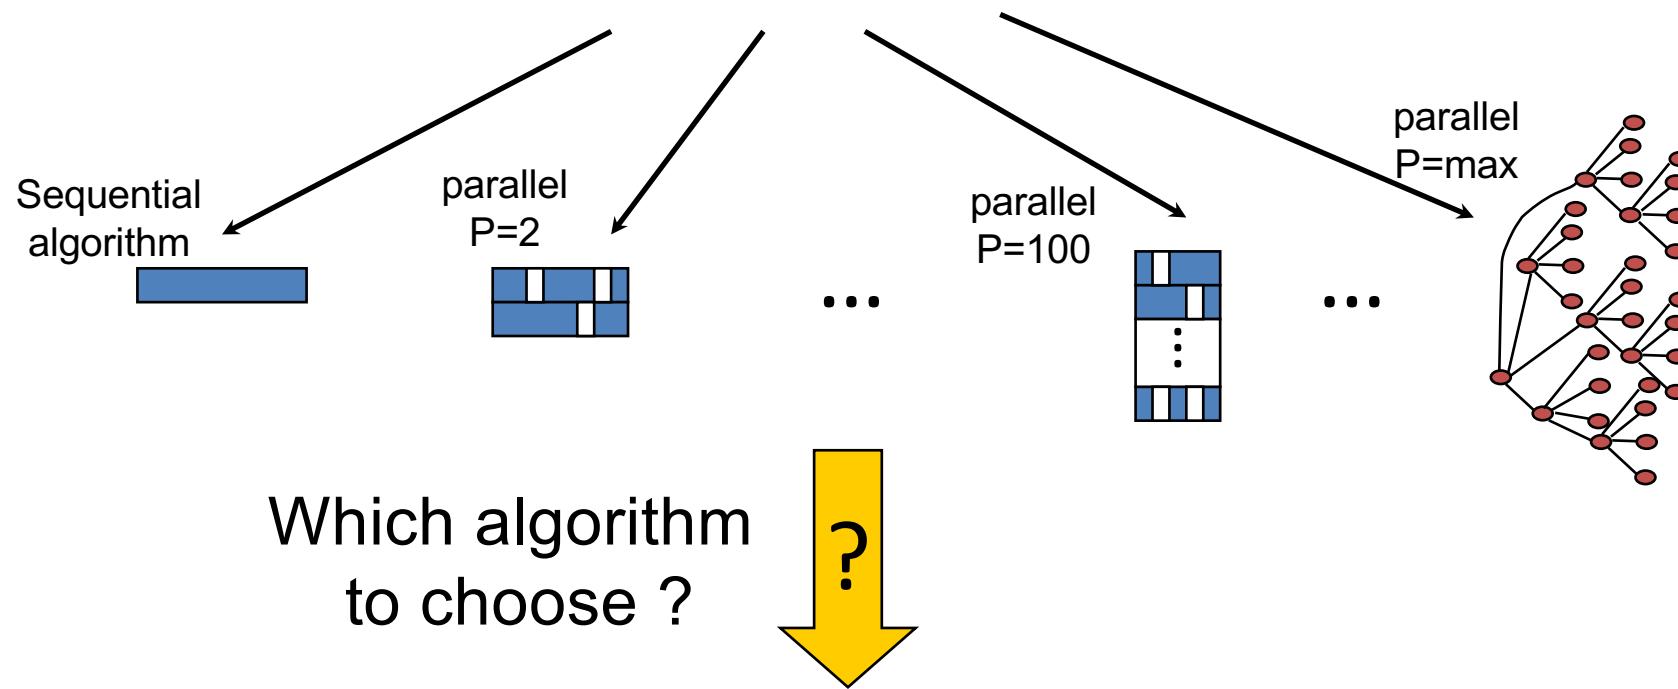
A kind of parallel algorithm, but heterogeneous “computing models”

Why interactive algorithms ?

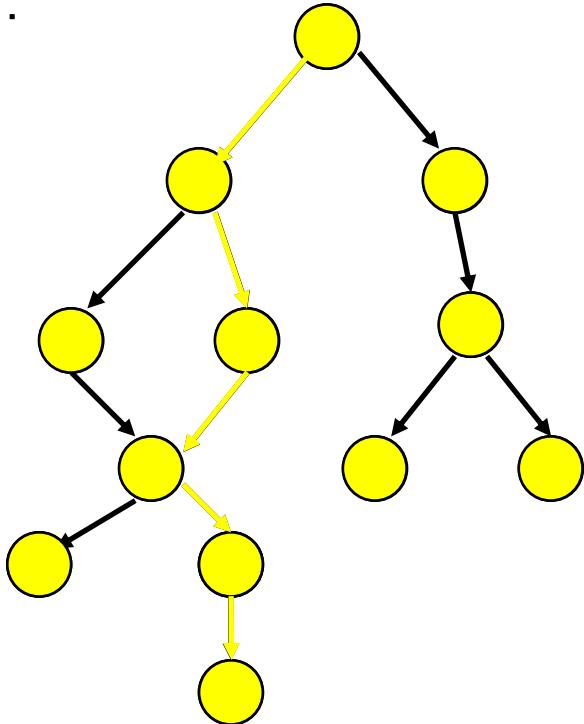
- For provable oblivious performances,
- For provable security (result correctness)

Towards oblivious algorithms

*To design a single efficient algorithm
with provable performances on an arbitrary architecture*



Basic notations: Work and depth



“Work” W = #total number operations performed

“Depth” D = #operations on a critical path
(~parallel “time” on ∞ resources)

Relation to execution time $T(p, \Pi)$

For any greedy *maximum utilization* schedule [Graham69, Brent70, Jaffe80, Bender-Rabin02]

$$\frac{1}{\Pi_{ave}} \text{Max}\left(\frac{W}{p}; D\right) \leq Time(p, \Pi) \leq \frac{1}{\Pi_{ave}} \left(\frac{W}{p} + D \right)$$

- There exist schedulers that reach the bound : $\frac{1}{\Pi_{ave}} \left(\frac{W}{p} + O(D) \right)$

Work-Stealing: a basis to design processor-oblivious algorithm

- **Work-stealing** = “*A decentralized thread scheduler: whenever a processor runs out of work, it steals work from a randomly chosen processor.*” [..., Leiserson 1998/Cilk, ...]
interaction between idle resources and busy ones
- Scheduler that reach: $\frac{1}{\Pi_{ave}} \left(\frac{W}{p} + O(D) \right)$
Moreover: if D small, few steals request [$O(p.D)$ w.h.p.]
- *When both W and D are small, leads to oblivious parallel algorithm*
 - *both theoretical and practical [Cilk, TBB, Kaapi, ...]*
 - *But, in general, W and D are antagonist ...*

Example: Parallel prefix

- **Prefix problem :**

- input : a_0, a_1, \dots, a_n
- output : π_1, \dots, π_n with

$$\pi_i = \prod_{k=0}^i a_k$$

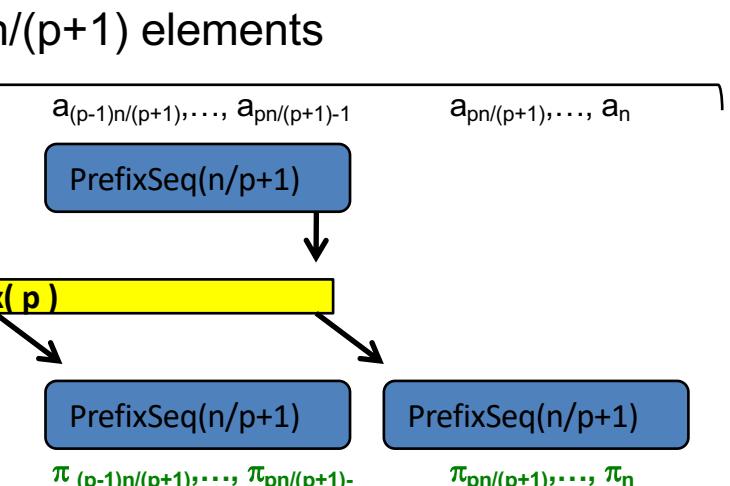
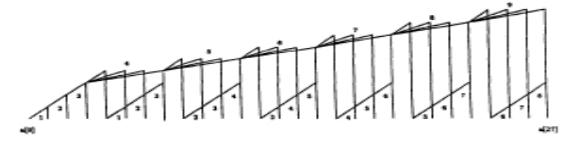
- **Sequential algorithm :**

- for ($\pi[0] = a[0]$, $i = 1 ; i \leq n; i++$) $\pi[i] = \pi[i-1] * a[i];$

*performs only **n** operations
(and minimal cache misses)*

- Optimal parallelization on **p identical processors**:

Optimal time $T_p = 2n / (p+1)$



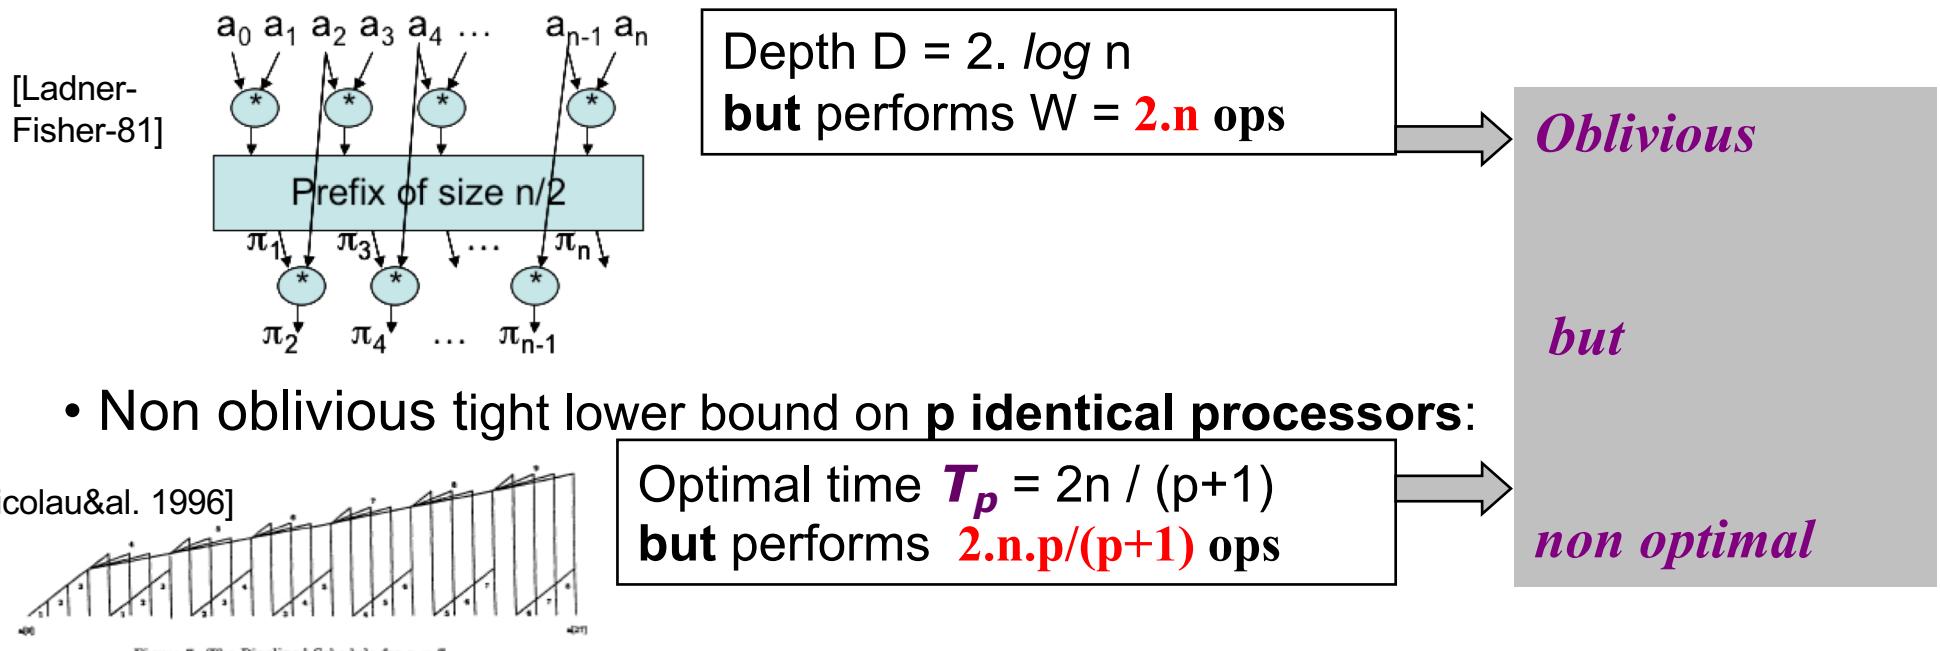
- **Non oblivious!**

Application to parallel prefix

- **Prefix problem :**
 - input : a_0, a_1, \dots, a_n
 - output : π_1, \dots, π_n with

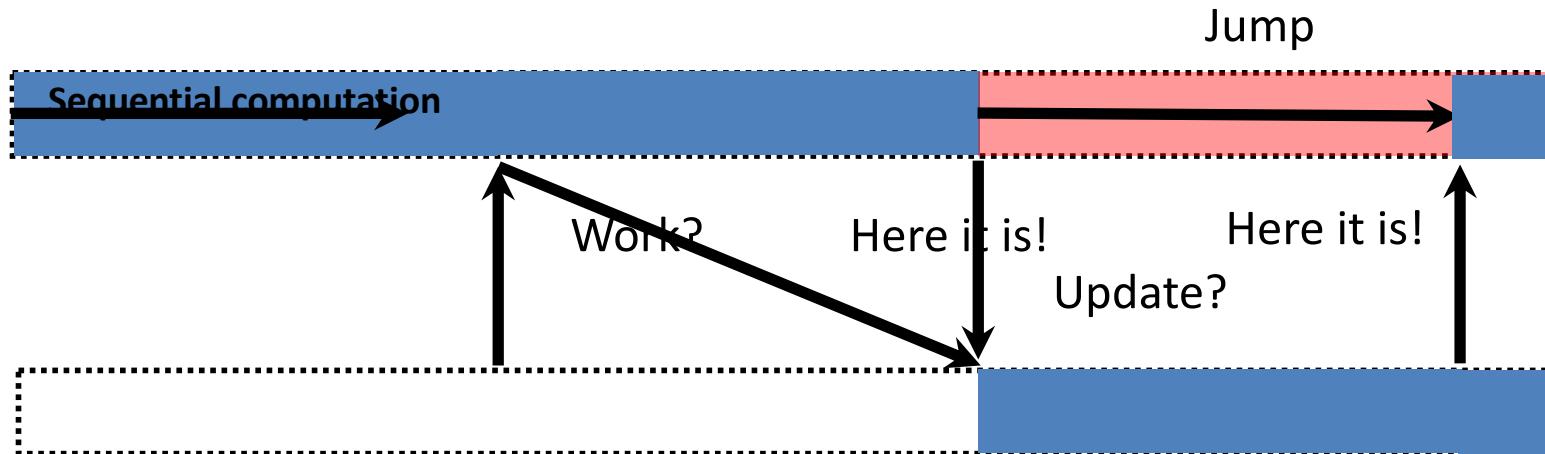
$$\pi_i = \prod_{k=0}^i a_k$$

- **Oblivious parallel algorithm** : recursive to minimize the depth D



Interactive parallel algorithm oblivious design

- Each resource performs a specialized sequential algorithm
 - other resources act as co-processors
 - at any time, sequential computation is in progress

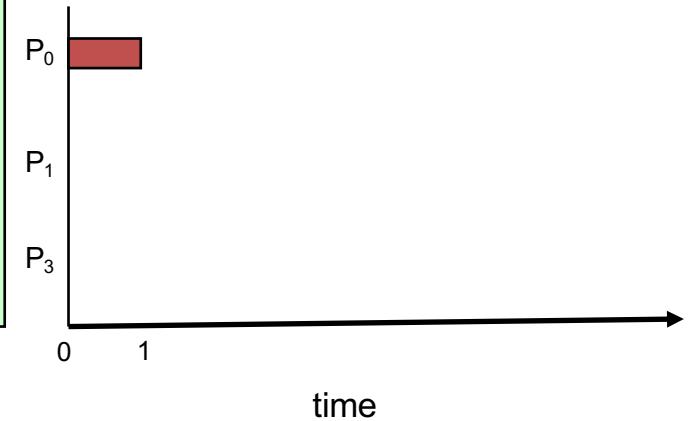
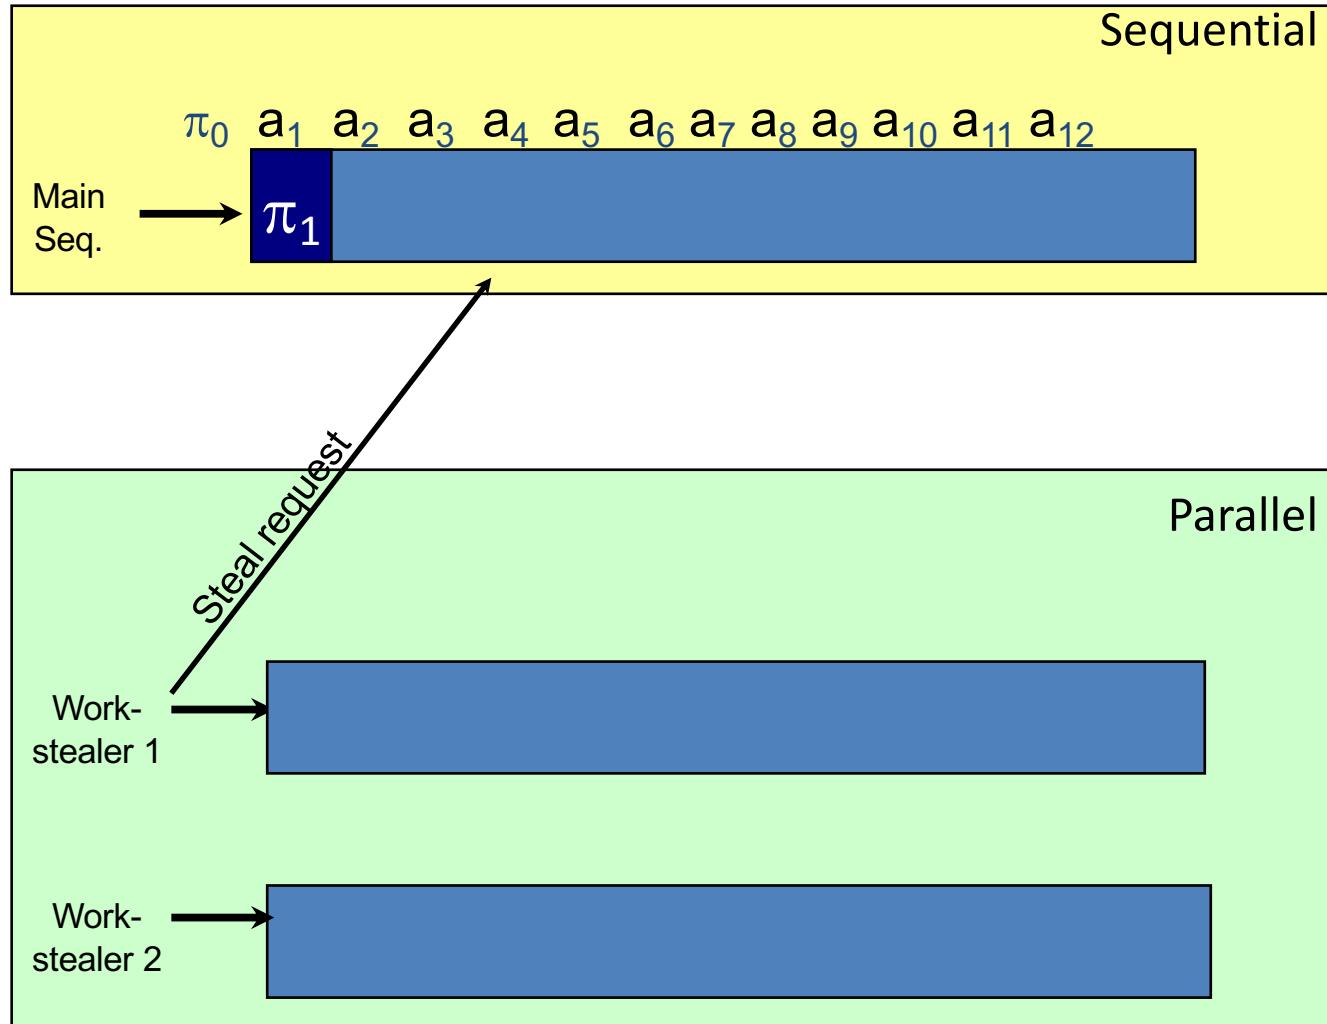


- When a stealer appears, it extracts some work (steal).
- When the work is completed, partial results are merged (eventually, preemption of the stealer).

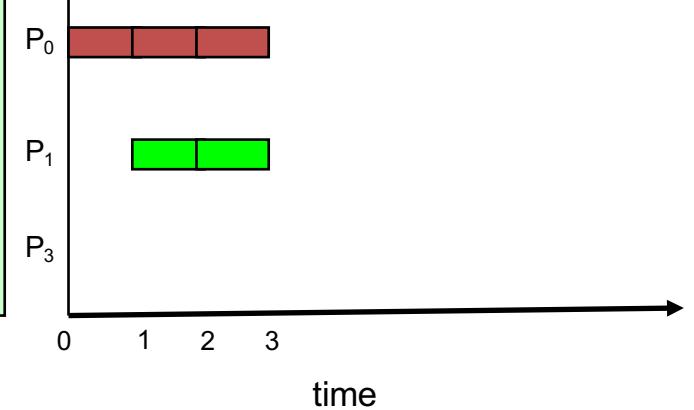
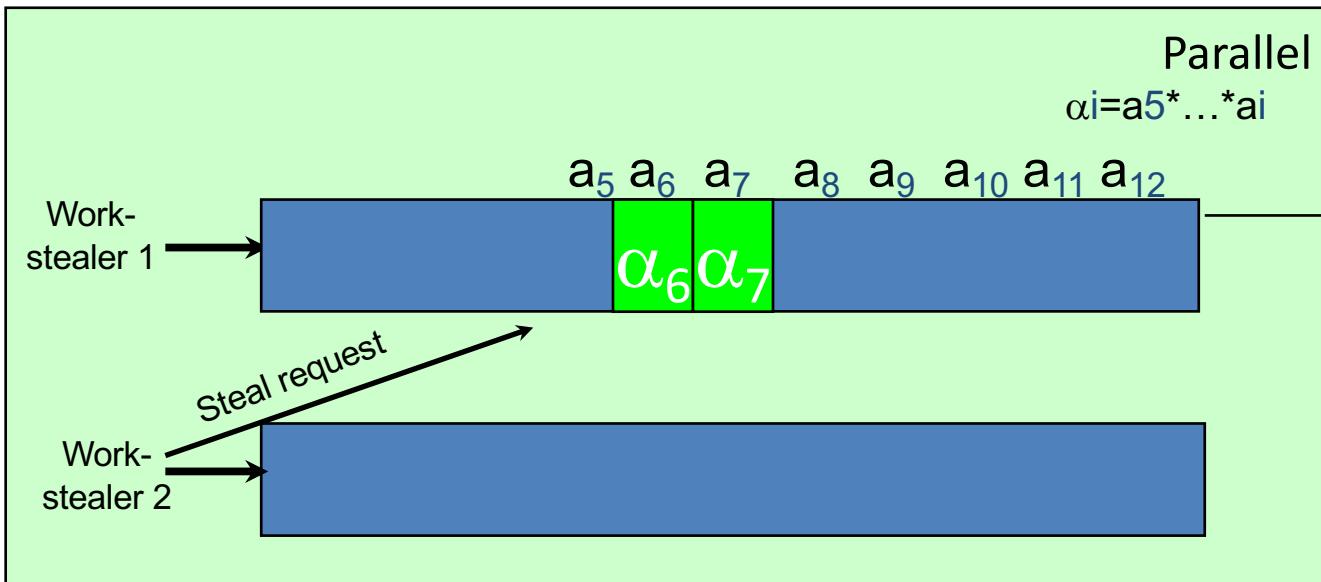
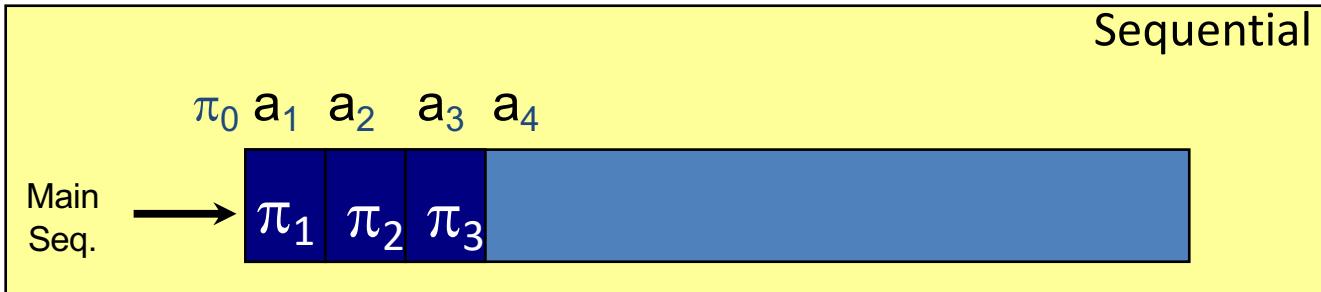
Provable performances in various contexts

- Computations with iterators: C++ STL algorithms
 - [Traoré, Roch, Maillard, Gautier, Bernard EUROPAR 08]
- Block Cipher Chaining modes (eg AES-CTR)
 - [Danjean, Gillard, Guelton, Roche, Roch PASCO 07]
- Anytime computation
 - [Soares, Raffin, Menier, Roch EGPV 07]
- Deterministic random number generator
 - [Mor, Roch, Maillard EUROPAR 14]

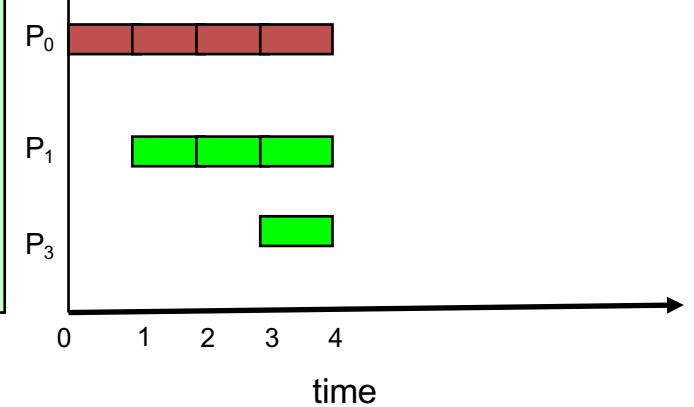
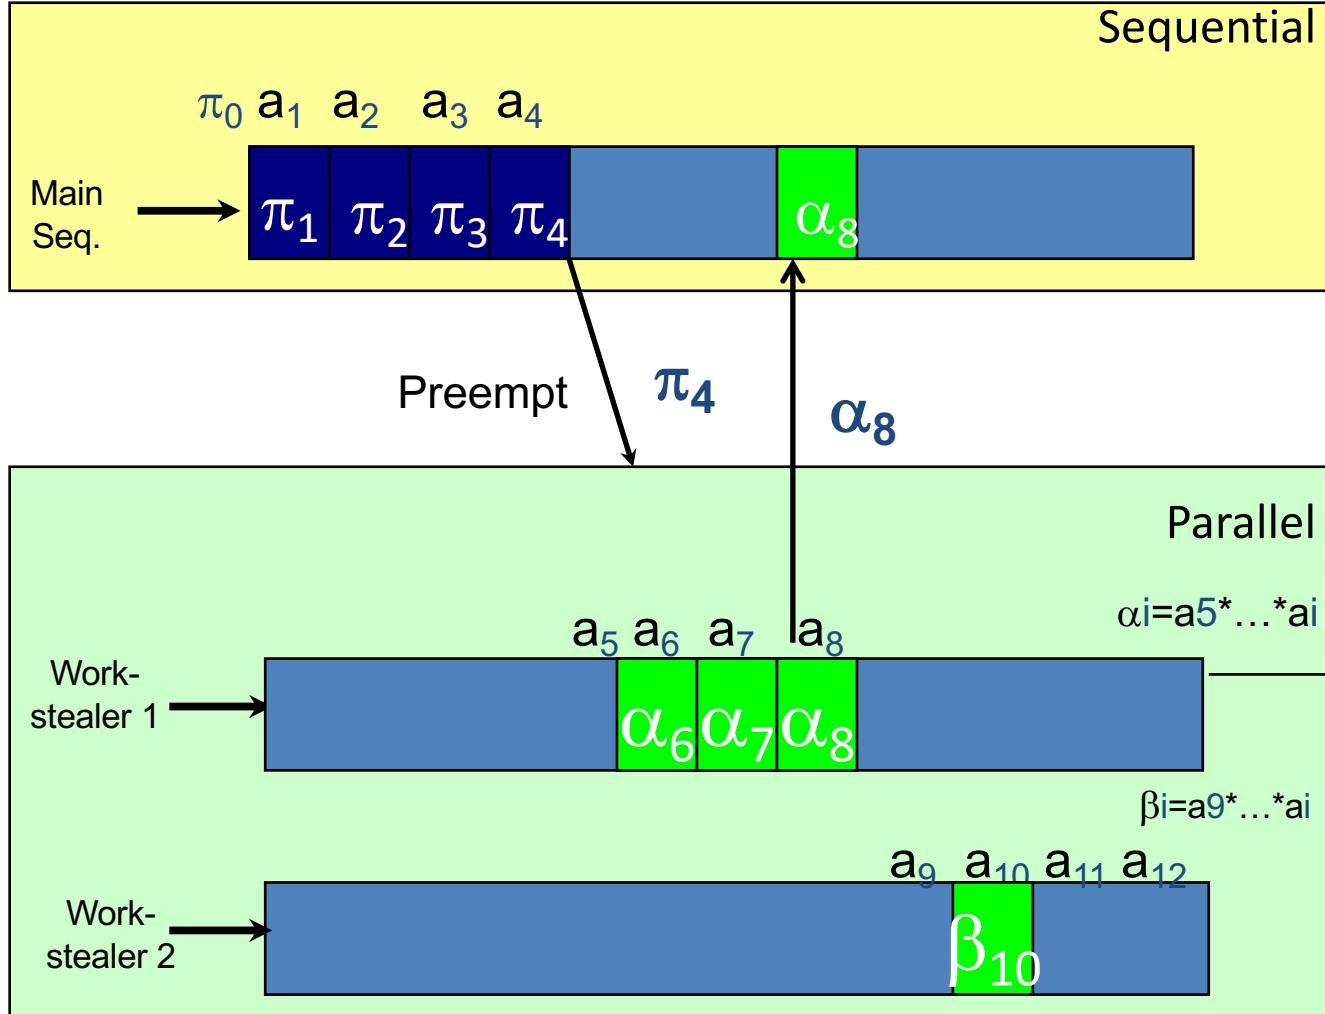
P-Oblivious Prefix on 3 proc.



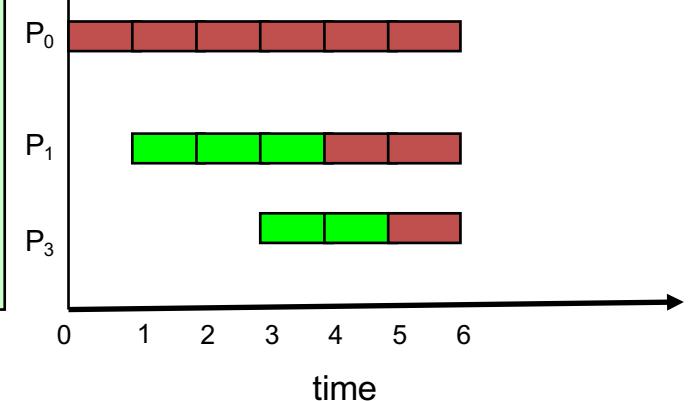
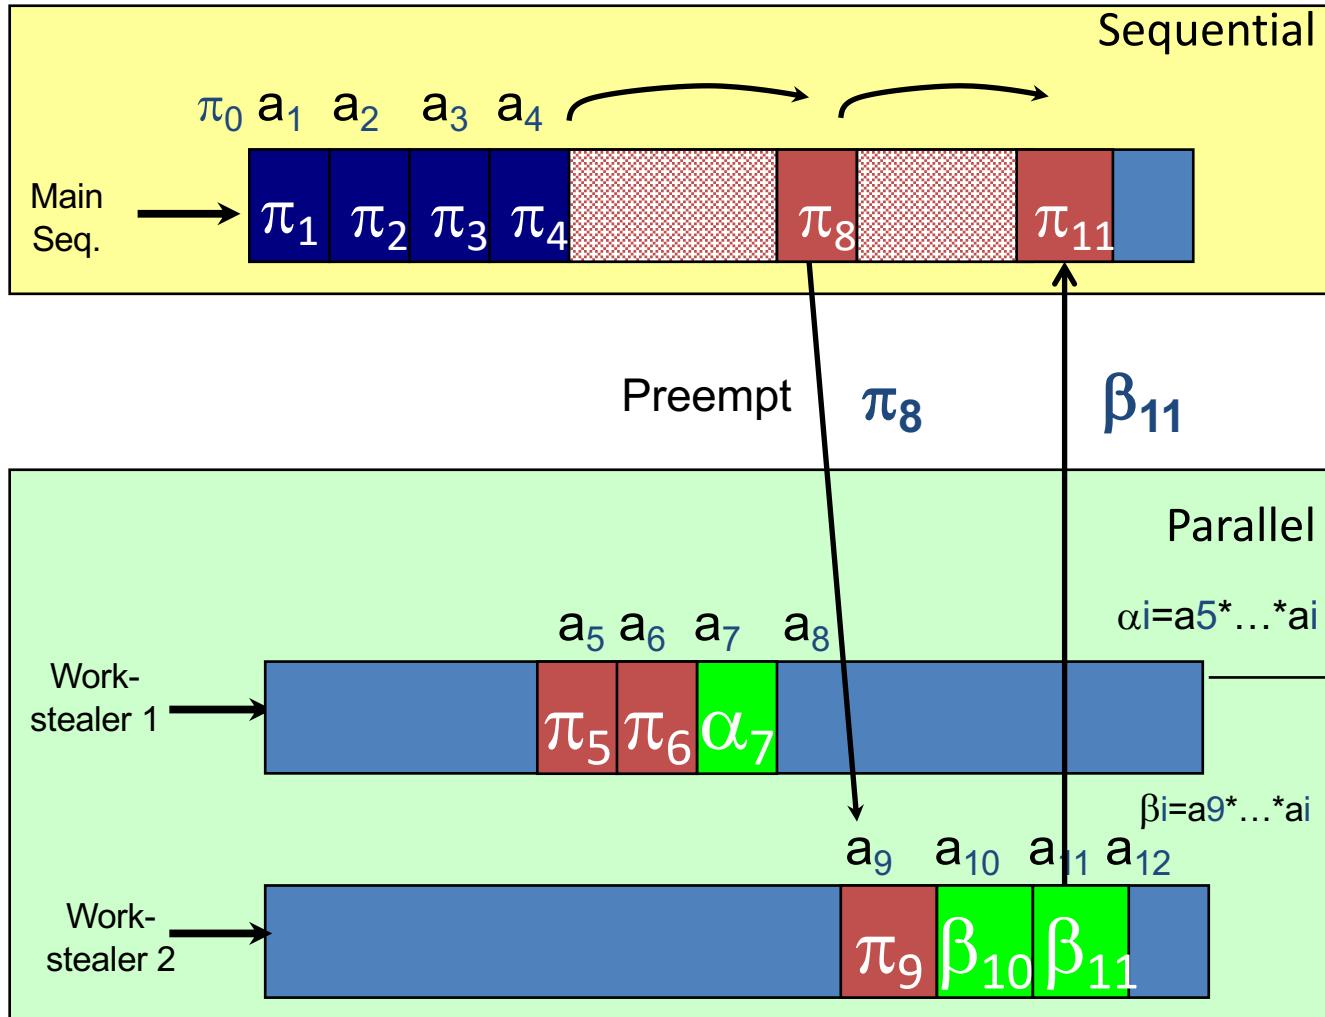
P-Oblivious Prefix on 3 proc.



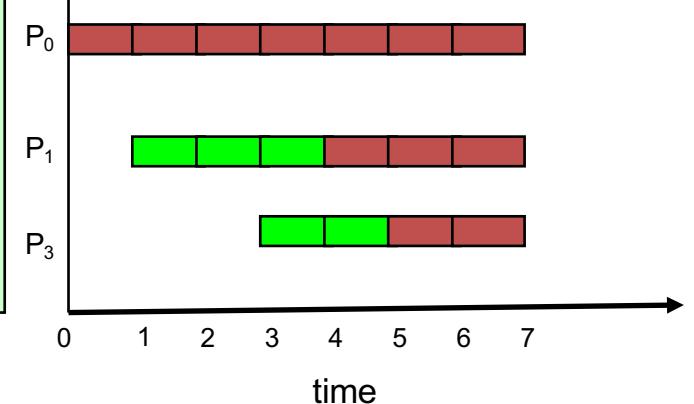
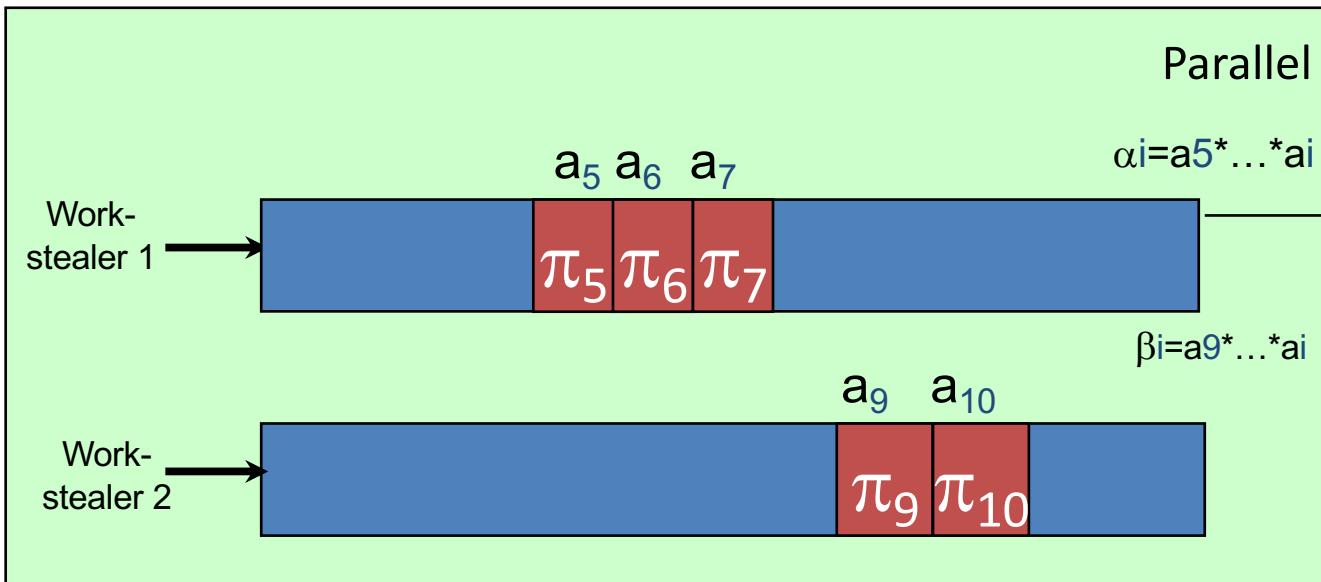
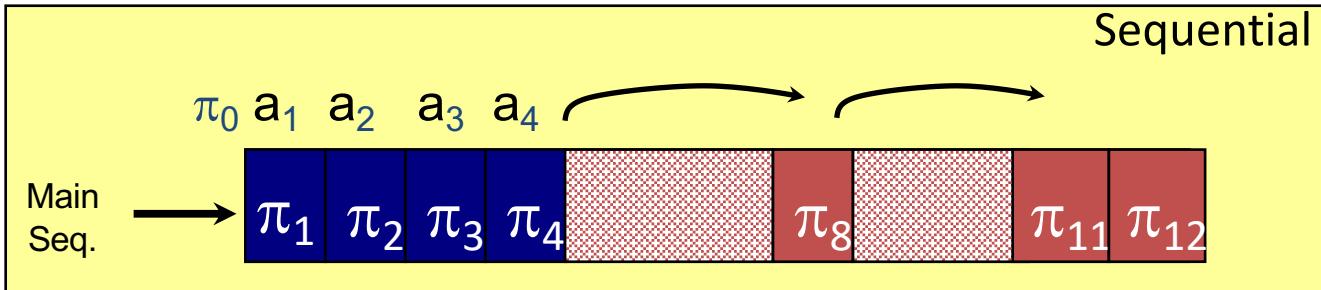
P-Oblivious Prefix on 3 proc.



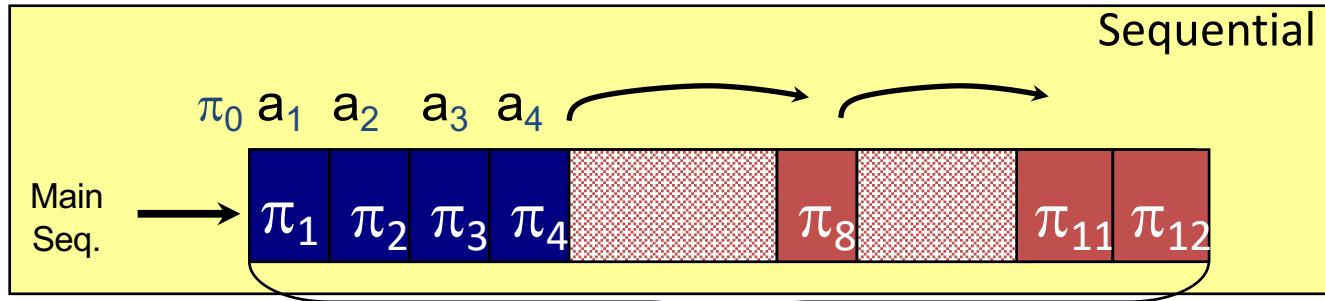
P-Oblivious Prefix on 3 proc.



P-Oblivious Prefix on 3 proc.



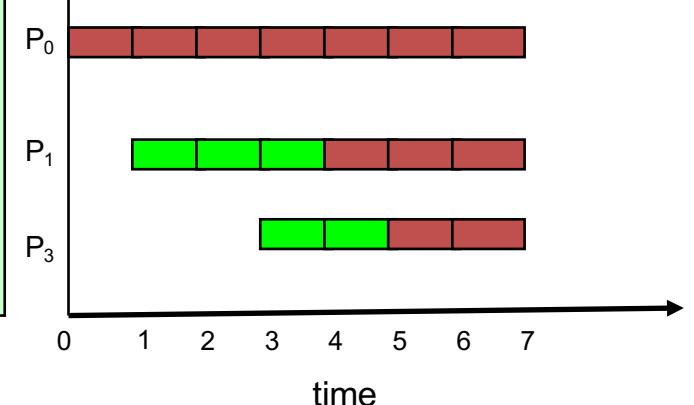
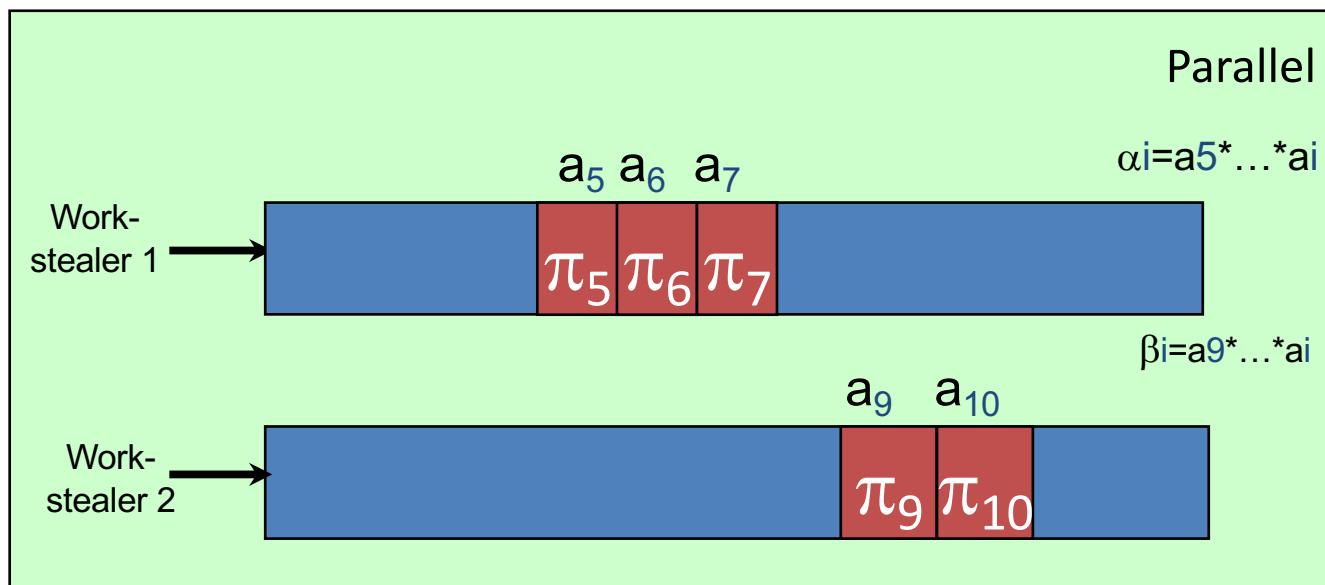
P-Oblivious Prefix on 3 proc.



Implicit critical path on the sequential process

$$T_p = 7$$

$$T_p^* = 6$$



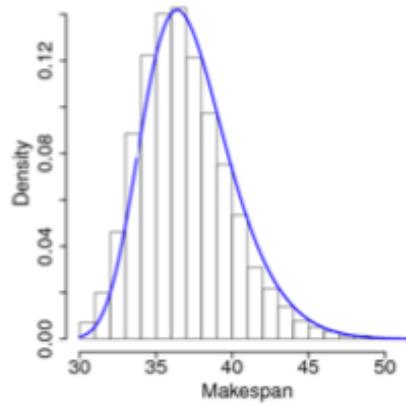
Provable performances

- Model parallel computation by its trace(s) (history)
 - Each computation is modelled a tree, which nodes are sequential blocks
 - #steals request = $O(p \cdot \text{Depth})$ w.h.p.
 - Non deterministic: family of all possible trees
- #Successfull steals, with « steal the oldest » strategy:
 - Let M = maximum number of forks on a parallel
 - Deterministic strategy: $\#\text{succ_steals} \leq M \cdot (p-1)$
 - Random steal strategy:
$$E[\#\text{succ_steals}] \leq (M-1)p \cdot \ln \ln p + p \cdot \ln p + O(P)$$
 - The finer the parallel grain => the fewer the tasks
=> the larger the effective granularity !!!

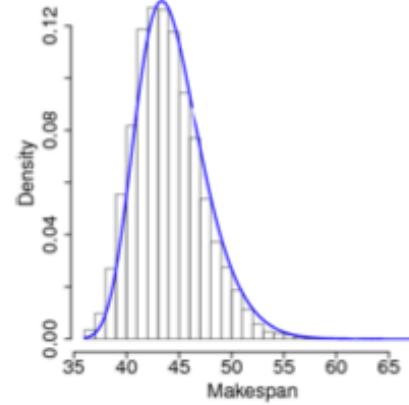
- Randomized Depth-First work-stealing : DAG with arity 2
- #average steals_requests per process = $O(T_\infty)$ [Arora&al 2001]
- Tighter analysis based (Distributed List Scheduling model) $< 5.5 T_\infty + O(1)$

[Tchiboukdjia, Trystram, Roch, Beranrd 2009, ISAAC 2010, A. Oper. Res. 2012]

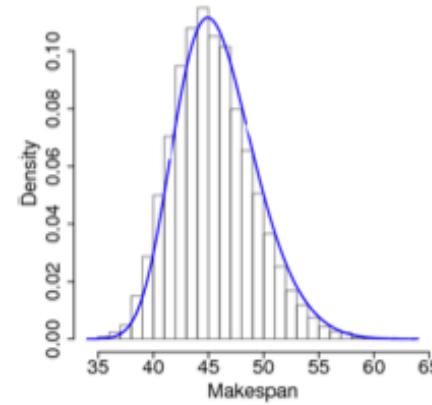
- independent unit tasks : # average steals_requests $< 3.65 \log_2 W + 3.64$
- with initial random allocation $< 1.83 \log_2 W + 3.64$
- with cooperative work stealing $< 2.88 \log_2 W + 3.4$



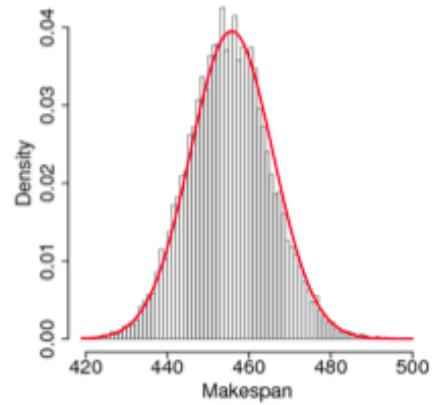
(a) Unit Tasks



(b) Weighted Tasks



(c) DAG (short D)



(d) DAG (long D)

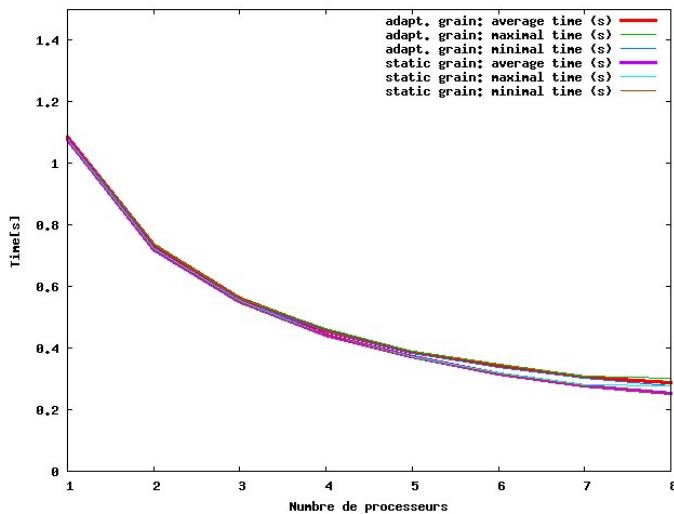
Provable performance of the Processor-oblivious parallel prefix

Execution time

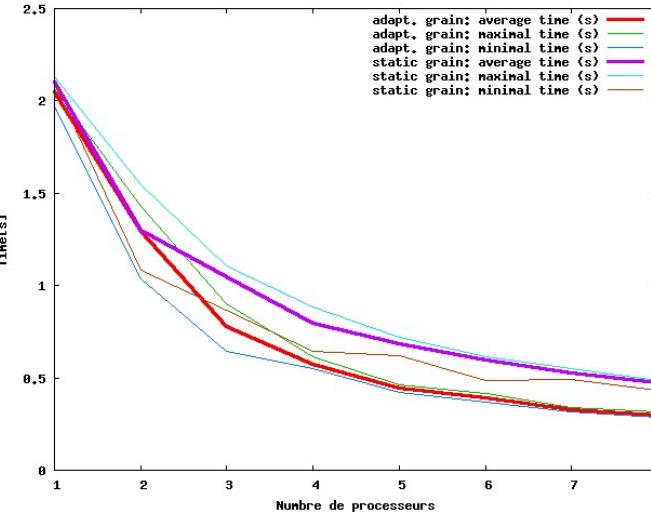
$$\leq \frac{2n}{(p+1) \cdot \Pi_{ave}} + O\left(\frac{\log^2 n}{\Pi_{ave}}\right)$$

Lower bound

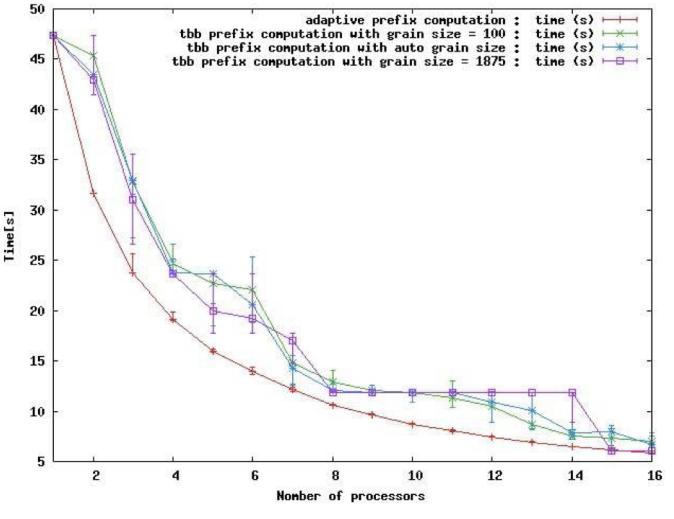
Practical performance [* = double, finest “grain” = 2048 double] [TRMGB, Europar’08]



Single user



Processors with small variations



Comparison with TBB

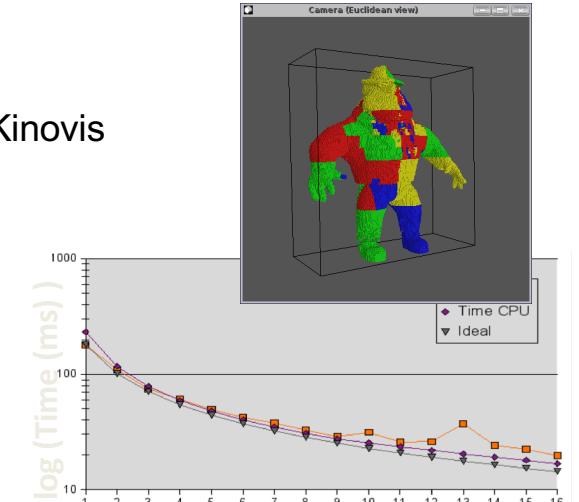
+ Cache-friendly (from the sequential algorithm)

Programming model

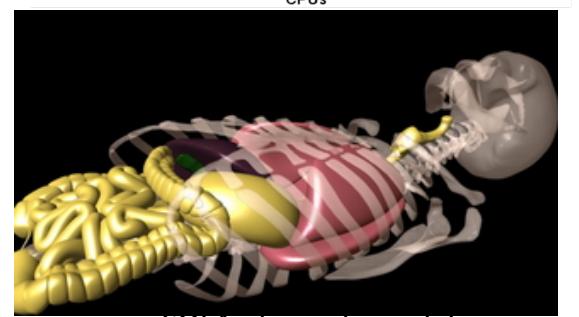
- Athapascan / Kaapi
 - Macro-data flow
 - Tasks with input-output dependencies
 - Dynamic task graph
 - Lazy computation of the DAG of tasks

Hybrid computing CPUs + GPUs

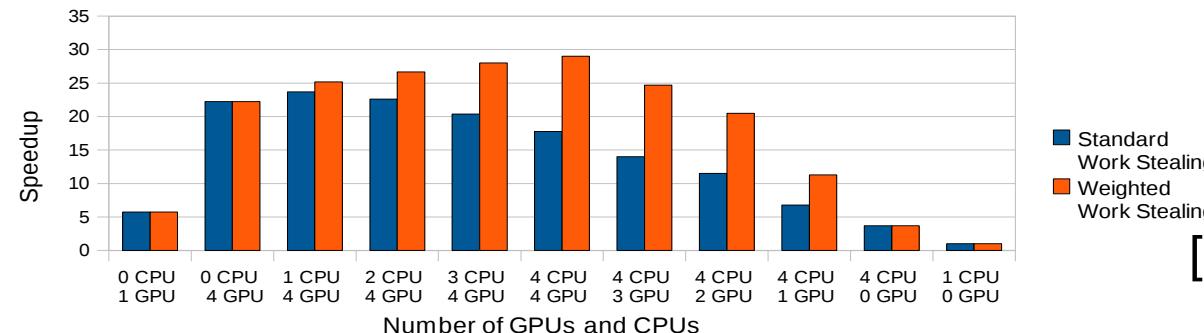
- parallel 3-D reconstruction [Grimage 2007, Kinovis 2015]
adaptive precision wrt frame rate



- Each task has 2 implementations:
CPU and GPU
 - Work-stealing + data-partitioning
KAAPI embedded in SOFA



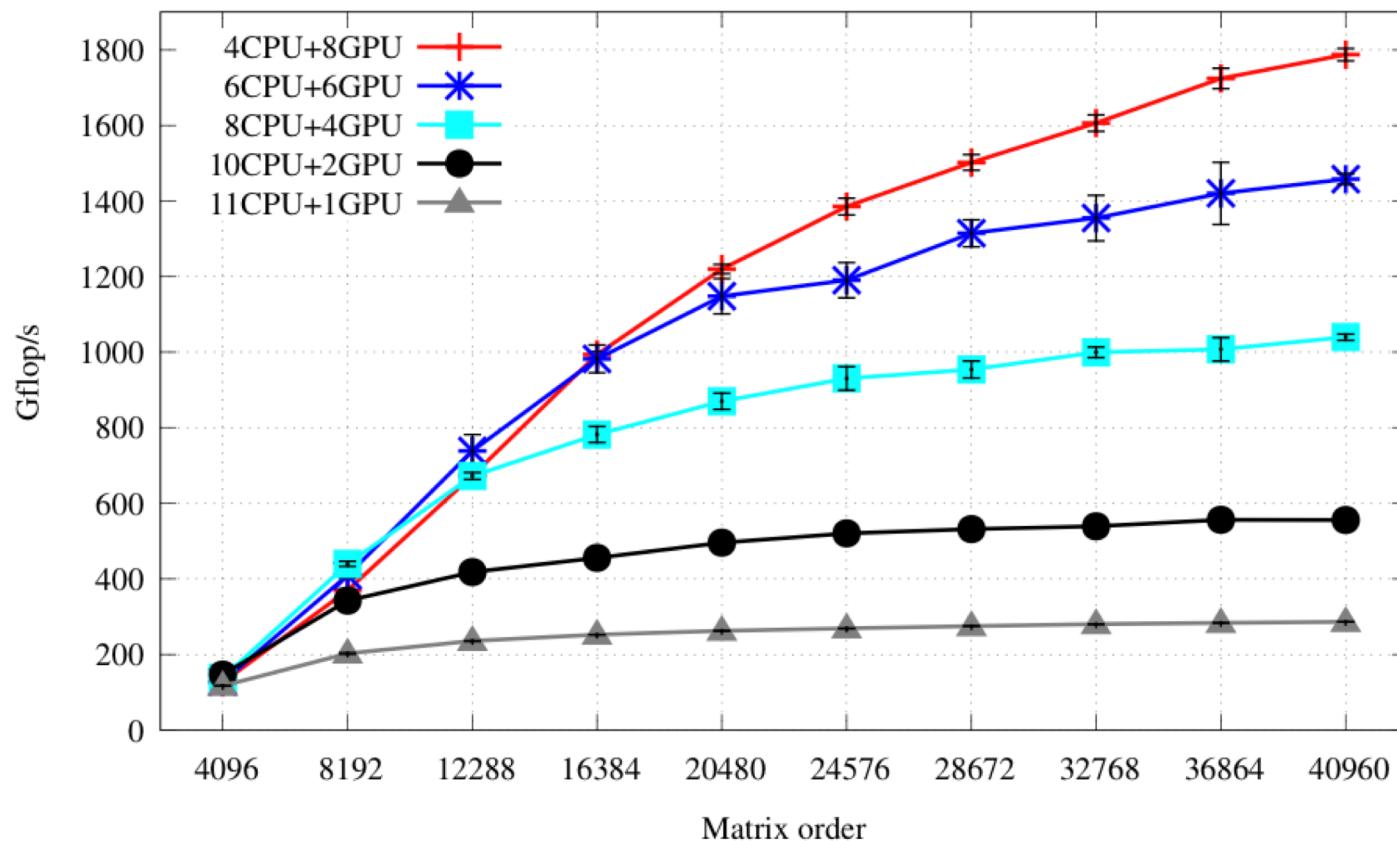
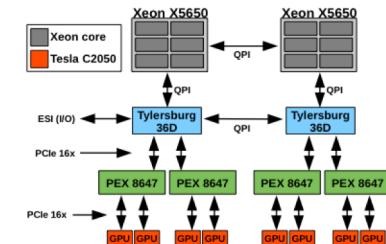
SOFA physical simulation
[Evasion / Imagine]



[E Hermann 2010]

DPOTRF multi-CPUs + Multi-GPUs

2 hexacores INTEL X5650 (2.66Ghz)
8 GPUs Tesla C2050 (fermi GPU, 1.15Ghz, 3GBytes)

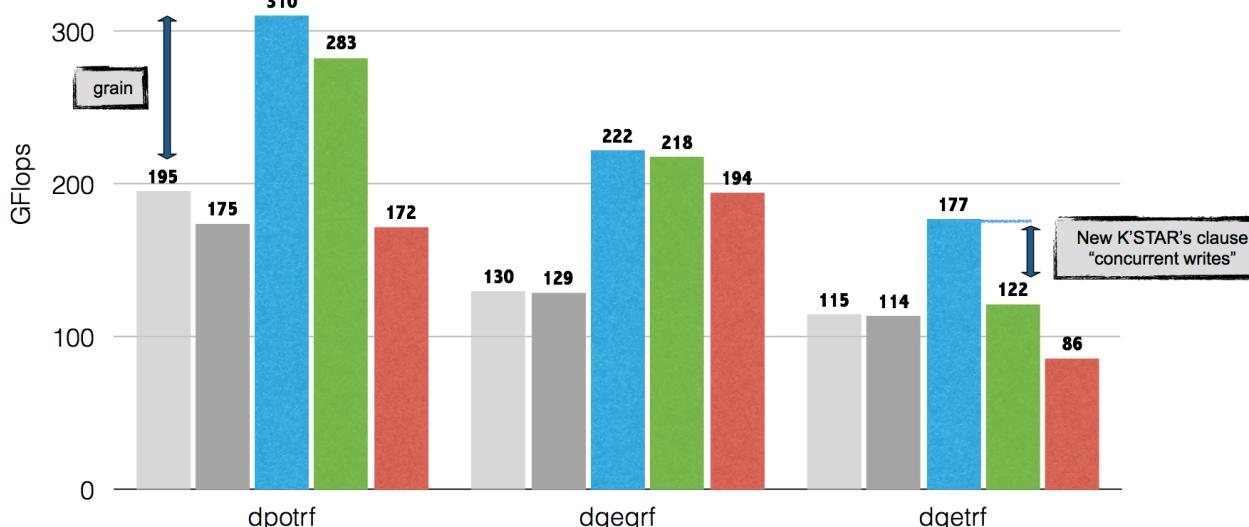


Gautier &al.
[IPDPS 2013]

Integration in OpenMP-4

KSTAR / GCC / OMPSS

plasma/static plasma/dyn kstar/xkaapi gcc/libgomp ompss



- KASTOR benchmarks HPAC matrix 4096 x 4096, BS=128 (+ IB=64 for QR)

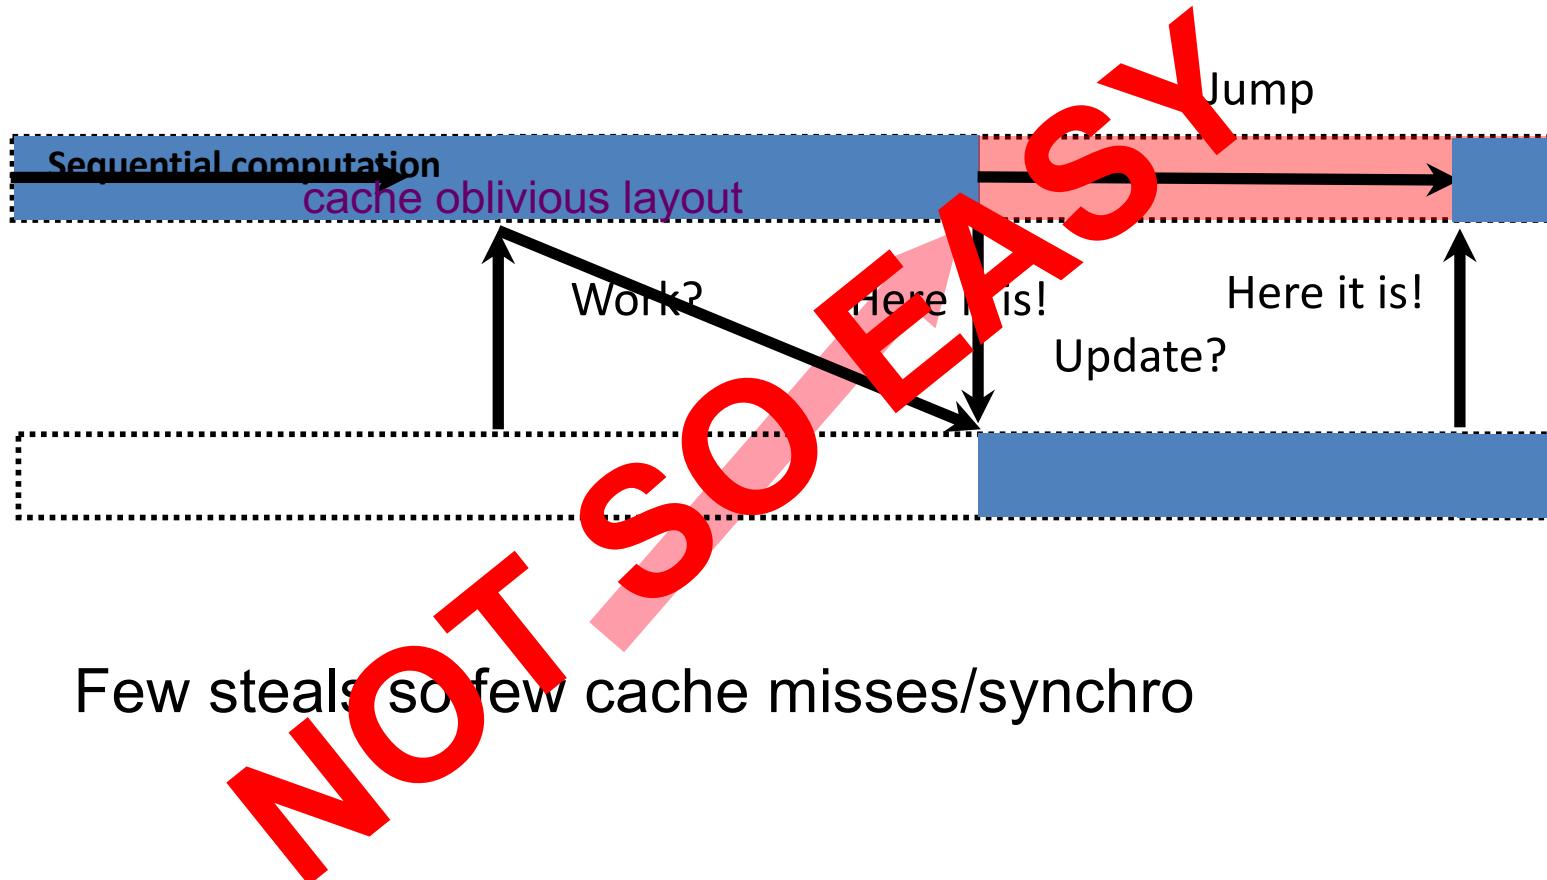


5

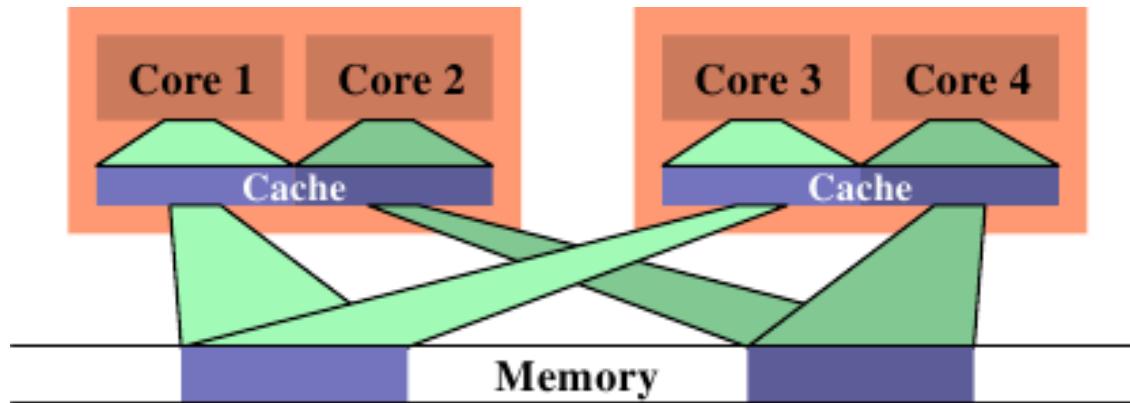
HPAC Machine
Intel Xeon CPU E5-4620 0 @ 2.20GHz
[Broquedis, Gautier &al. IWOMP 2014]

- 4 NUMA nodes
- 32 coeurs

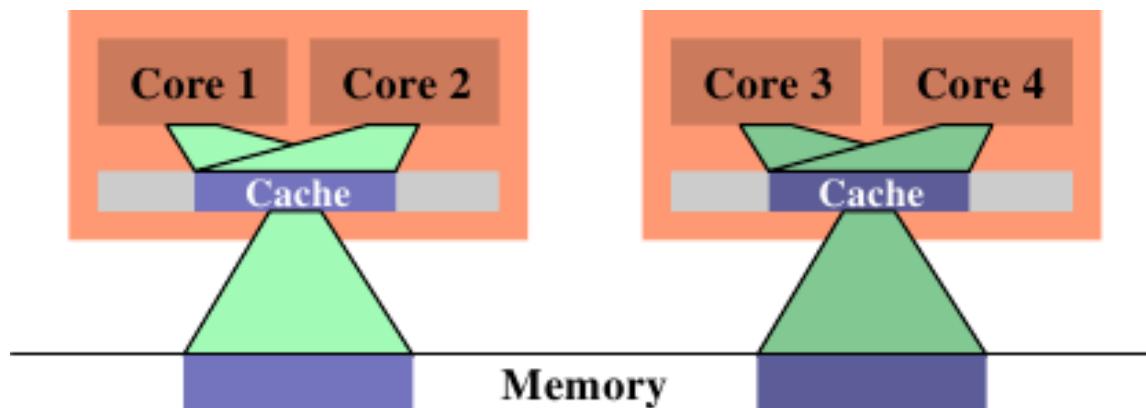
What about locality?



Interaction cache / scheduling



Can be worse than
each core has a cache of half size



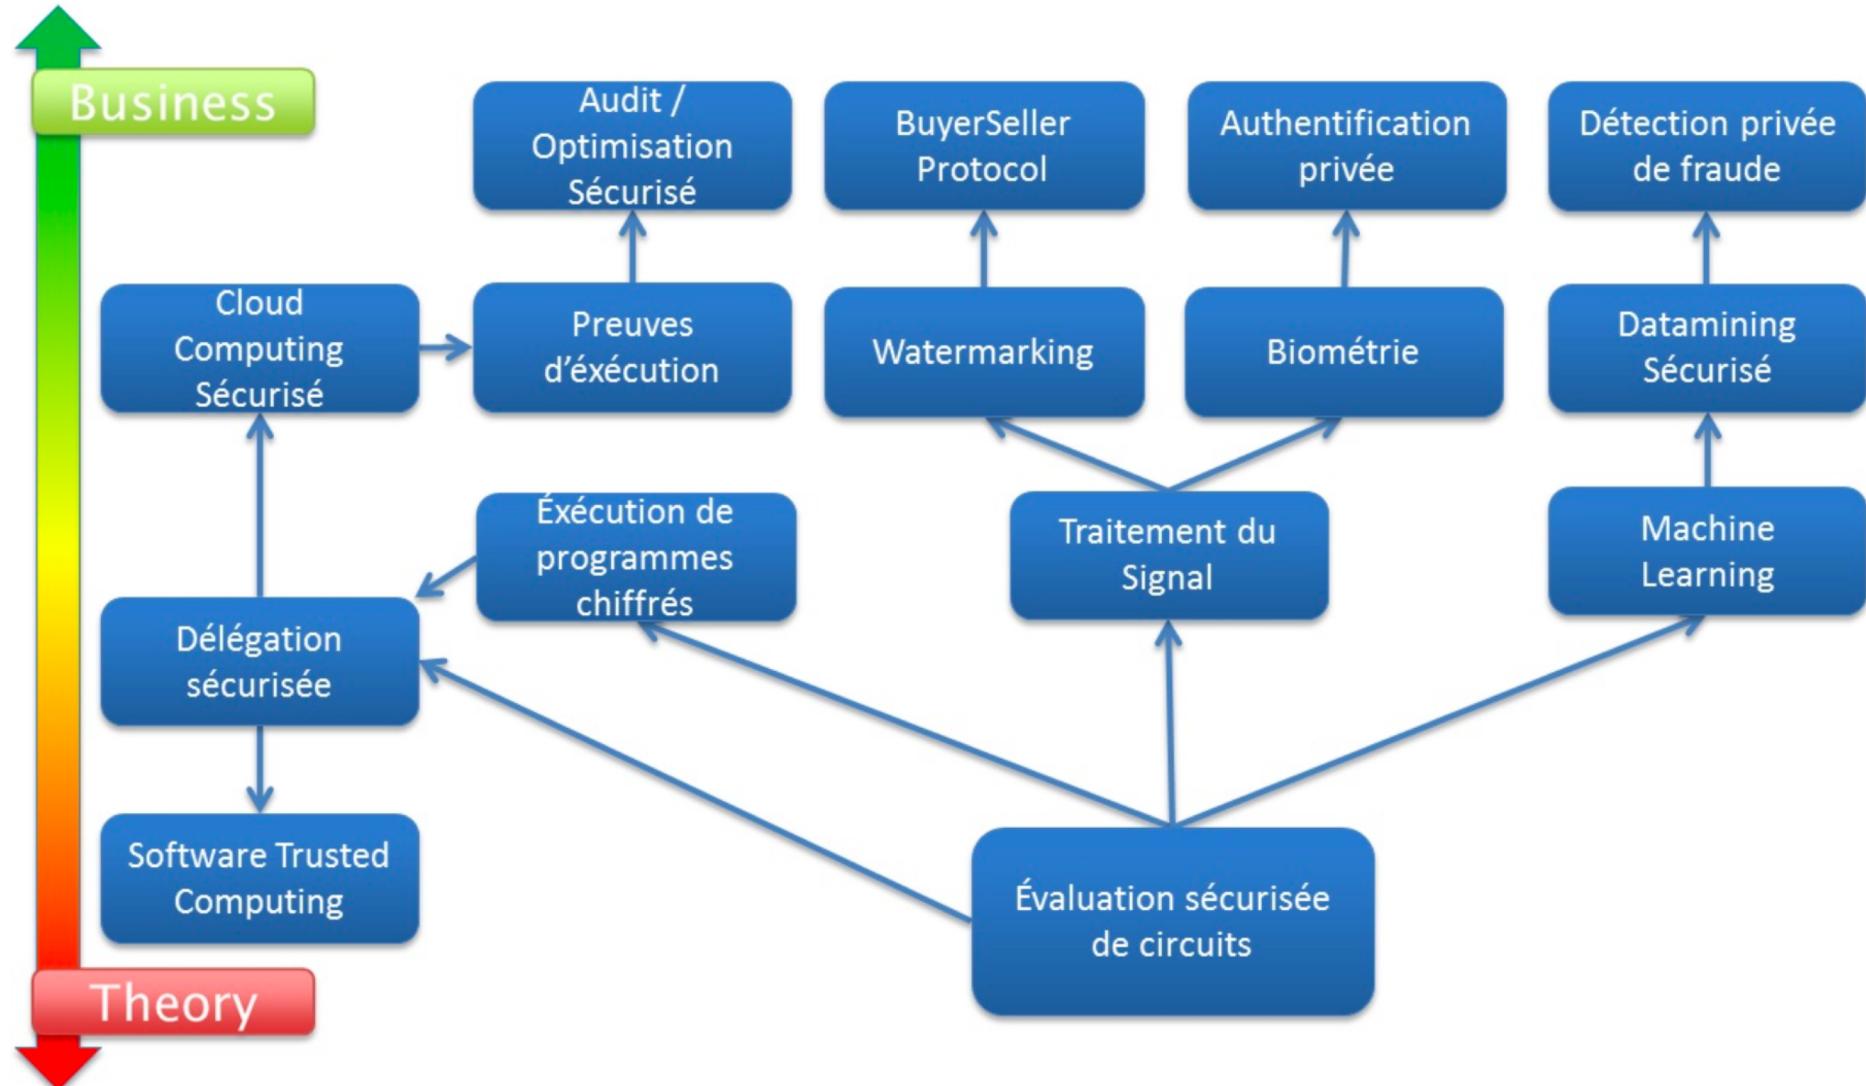
Cache sharing

Can be as good as
each core has a cache of full size

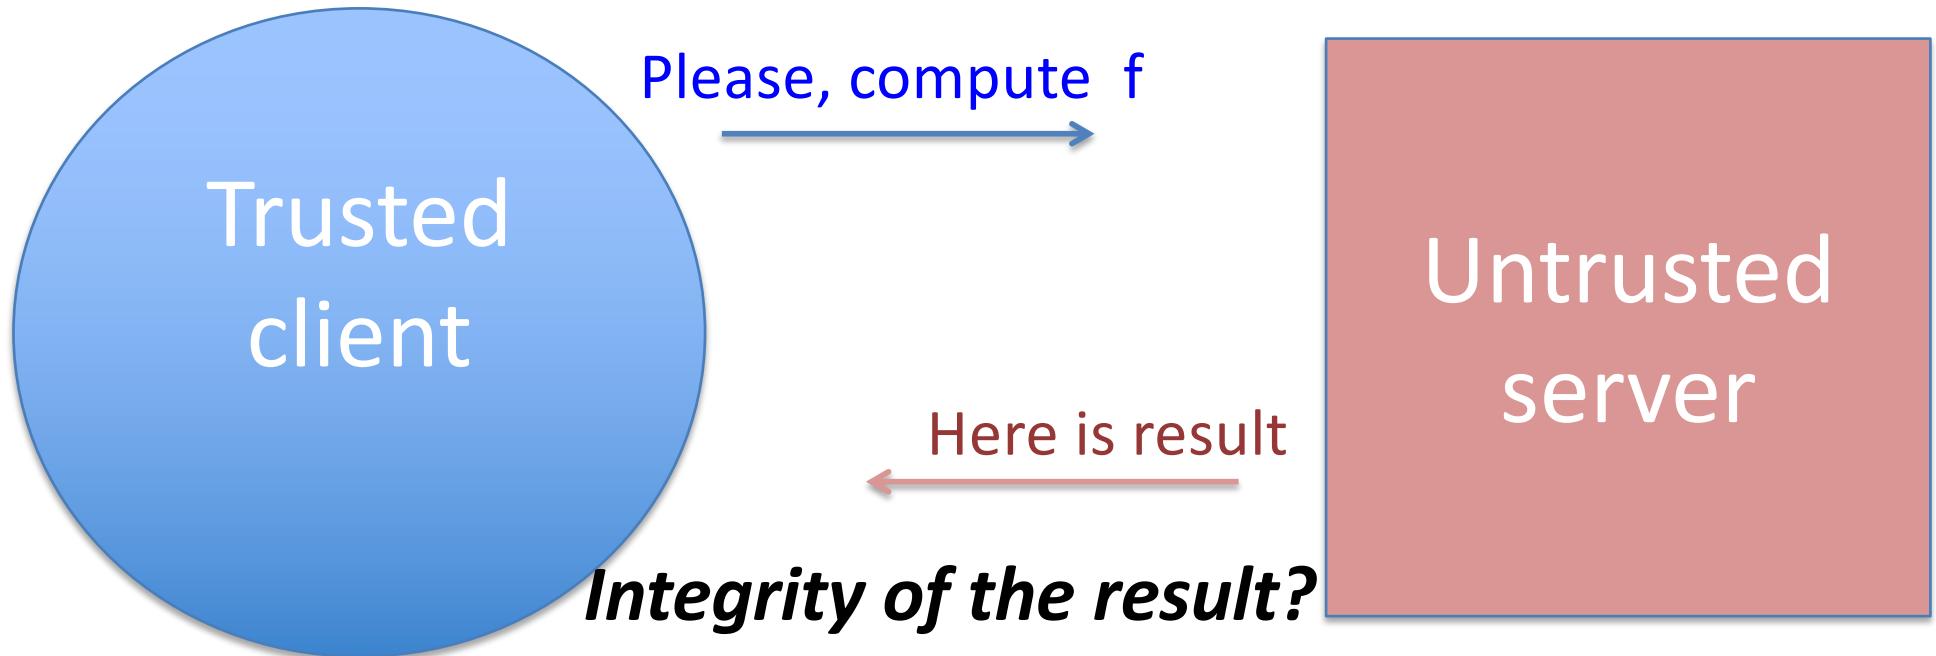
Plan

Interactive computations and outsourcing

1. For oblivious parallel provable performances
 - Work-stealing
 - Hybrid computing
2. For provable security
 - Result correctness (integrity)



Delegating computation



- Contexts
 - Co-processor (overclocked...)
 - Supercomputer (soft errors)
 - Cloud computing
 - Volunteer computing

Motivating example

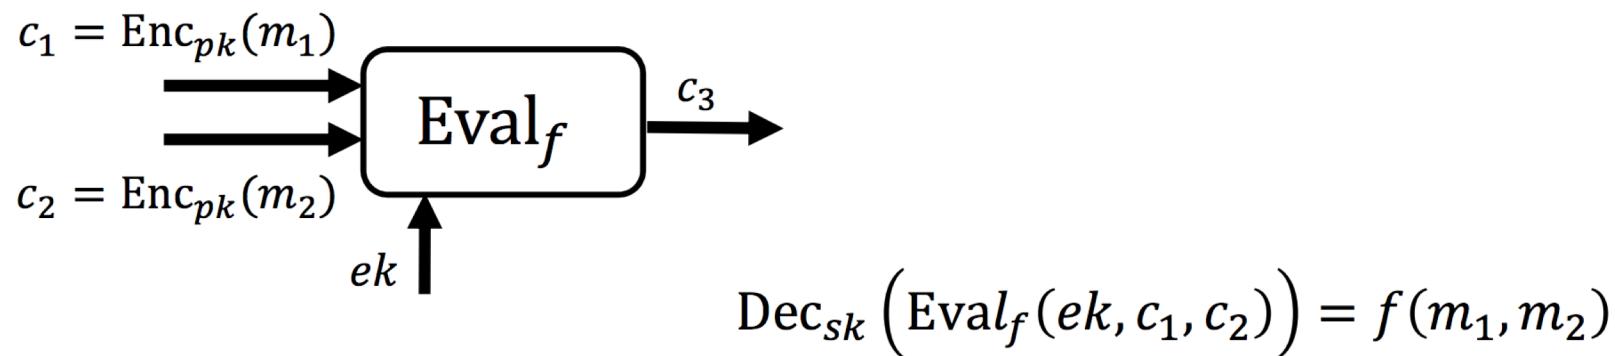
- Peggy has developed a nice application that efficiently solves Traveling Salesman Problem
- Victor sends Peggy the location (map) of his clients and pays her for the shortest Hamiltonian circuit
- How does Victor check he gets the shortest?

Outsourcing and privacy

- Previous example: Victor provides the location of his clients => how to keep this secret
- Other example: digital doctor

Outsourcing and privacy

- Outsourcing computation with secret input
 - Computation is performed on encrypted data
 - Based on asymmetric encryption (eg thanks to **fully homomorphic encryption**) [Gentry 2009]



Homomorphic encryption: El Gamal e-vote

- **Remind El Gamal** (in cyclic group G with g a generator):
 - Bob has private key b and public key $B=g^b$
 - Alice: $C=E(M) = (c_1=g^r, c_2=M \cdot B^r)$ Bob= $D(C)=c_1^{-b} \cdot c_2 = M$
- **El Gamal enables homomorphic addition** (*note Alice encrypts g^M instead of M*):
 - $C=E(M) = (g^r, g^M \cdot B^r)$ (encode g^M instead of M)
 - $C'=E(M')=(g^{r'}, g^{M'} \cdot B^{r'})$
 - Multiplication of ciphertext in G matches addition of plaintext (*e.g. integers plaintext*)
$$C \cdot C' = (g^r \cdot g^{r'}, g^M \cdot B^r \cdot g^{M'} \cdot B^{r'}) = (g^{r+r'}, g^{M+M'} \cdot B^{r+r'})$$

$$= E(M+M')$$
 - Enables anyone to compute as many additions of ciphertexts as desired
- **Application: electronic vote by homomorphic addition**
 - Each voter (Alice) sends her encrypted vote v (0 or 1) to the voting machine (Bob) :
 - $C(0)=(g^r, B^r)$ $C(1)=(g^{r'}, g \cdot B^{r'})$: each Alice checks her encrypted vote is correctly stored
 - Each one can compute the encrypted score of the vote : $\prod_{C \text{ voter}}(C) = (g^{\sum r}, g^{\sum v} \cdot B^{\sum r})$
 - The voting machine knows secret b : it decrypts the score $\sum_{\text{voters}} v$ and publishes it !
 - Why the voting machine does not need to compute discrete log ?

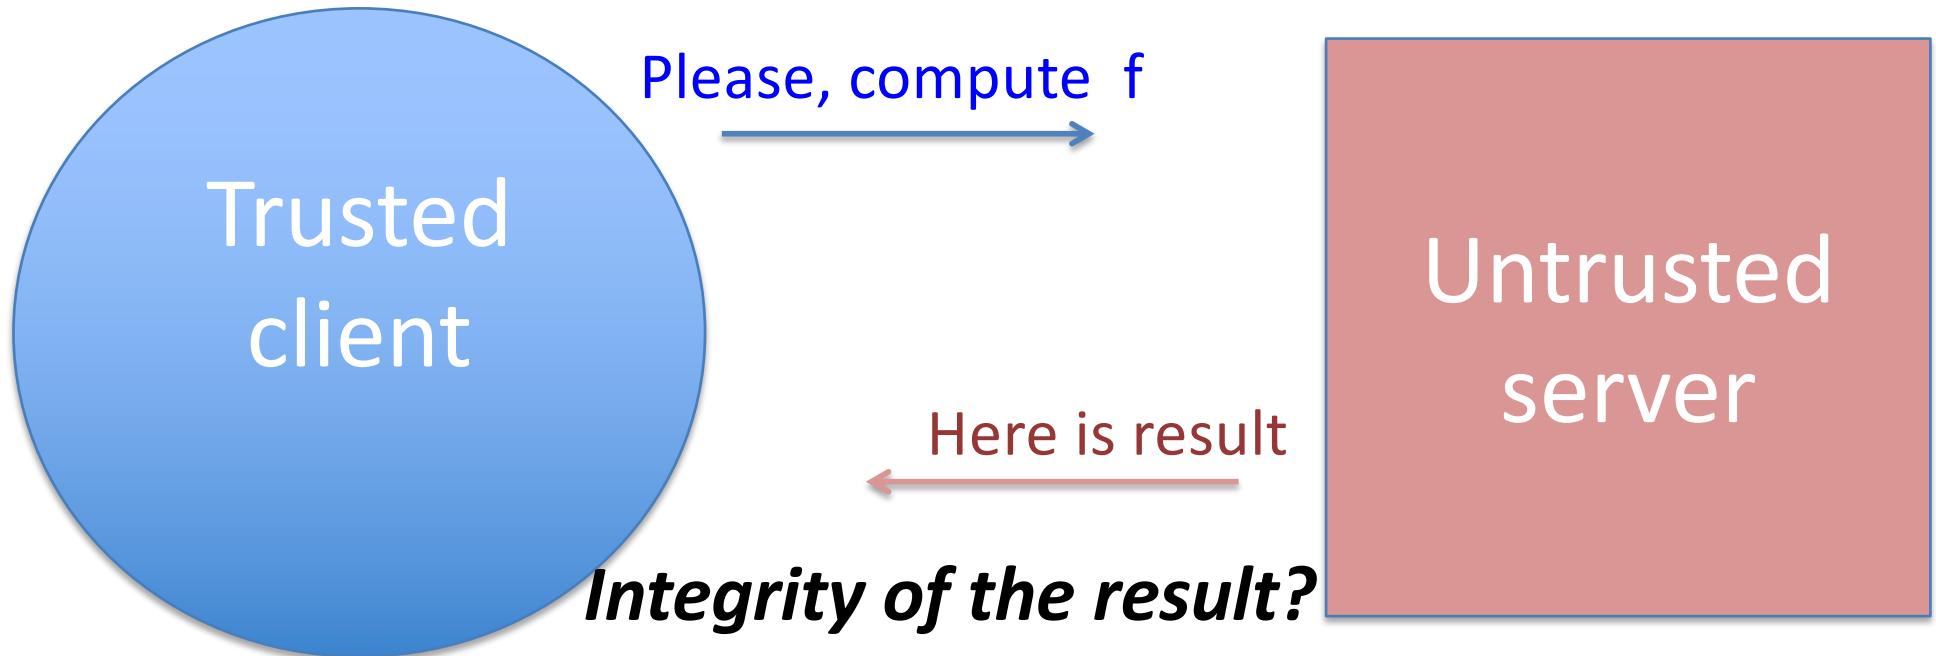
Fully homomorphic encryption

- How to perform AND and XOR on ciphered booleans ?
 - [Marten van Dijk, Craig Gentry, Shai Halevi, Vinod Vaikuntanathan]
 - Secret p : a large odd integer (eg thousands of digits)
 - For $x \in \{0,1\}$: $E(x) = pq + 2.r + x$
with random $q \sim$ million of digits and $r \sim$ twenty digits (the *noise*)
 - Knowing p : $(E(x) \bmod p) \bmod 2 = (2.r+x) \bmod 2 = x$
 - Without knowing p : $E(x)$ seems to give no information
- Fully homomorphic with $@x$ and y booleans:
 - $E(x)+E(x')=p.(q+q')+2(r+r')+x+x' \Rightarrow \bmod p \bmod 2 = x \text{ XOR } x'$
 - $E(x).E(x')=p(pqq'+q(2r'+x')+q'(2r+x)) + 2(2rr'+rx'+r'x) + x.x' \Rightarrow \bmod p \bmod 2 = x \text{ AND } x'$
- But with operations, the noise increases may become larger than p
- Key Gentry's idea: bootstrapping
 - Eliminate the noise by outsourcing remainder by p on the cipher domain
 - by homomorphic computation with AND and XOR (so without revealing p , only its ciphering !)
 - Care : the number of AND and XOR keeps the noise bounded
 - choice of p and the q 's large enough!

Outsourcing and privacy

- Homomorphic scheme enables to outsource encryption with secret key (or signature)
- Homomorphic encryption enables publicly Verifiable computation [Fiore, Gennaro 2012, ...]
 - Server computes on private data and produces a verifiable digest of the computation
 - Verification of the computation
 - Different from a result certification

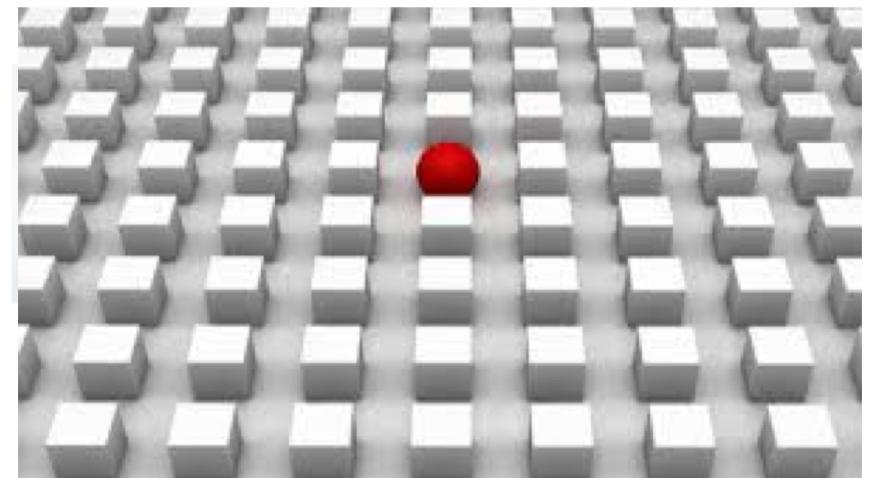
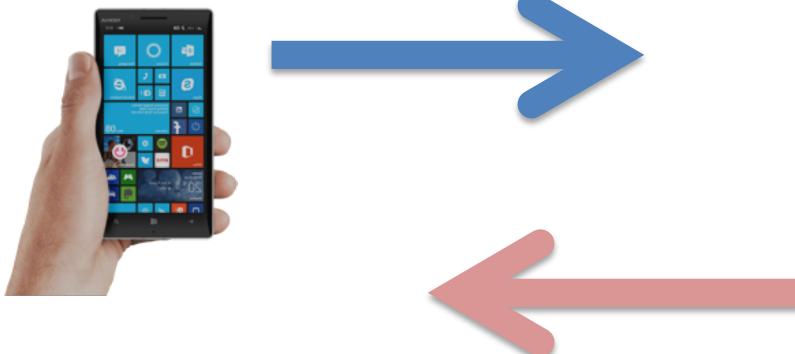
Delegating computation



- Contexts
 - Co-processor (overclocked...)
 - Supercomputer (soft errors)
 - Cloud computing
 - Volunteer computing

Attack models

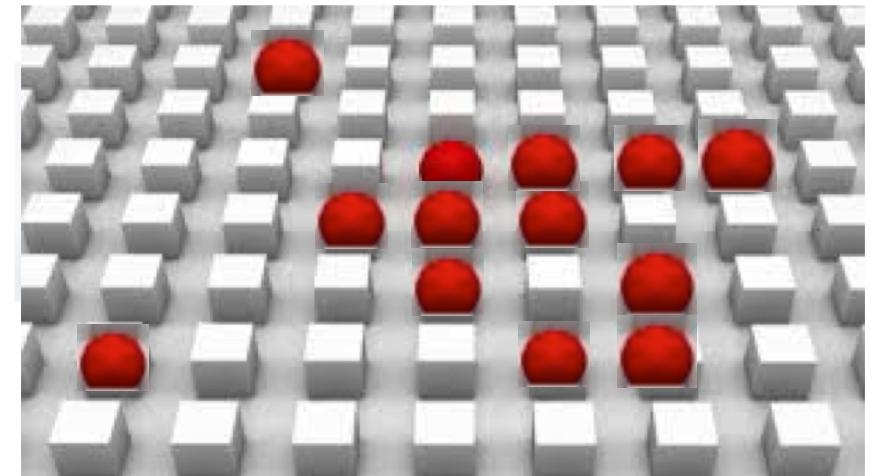
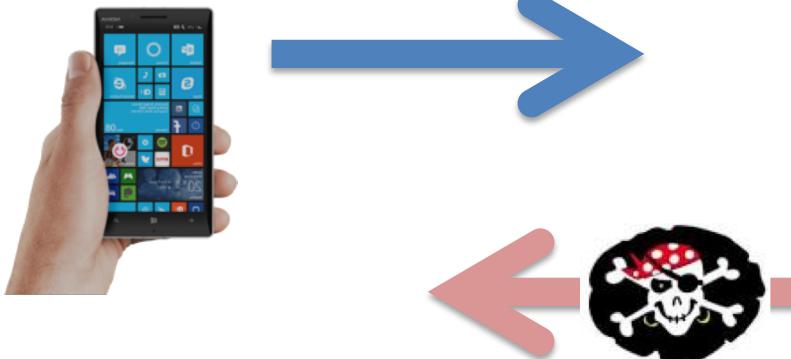
- No attack [current HPC and grid computing platform]
 - Failure (MTBF)



- **Attack on few isolated resources**
 - Soft errors - corruption of part of the computation

Attack models

- No attack [current HPC and grid computing platform]
 - Failure (MTBF)

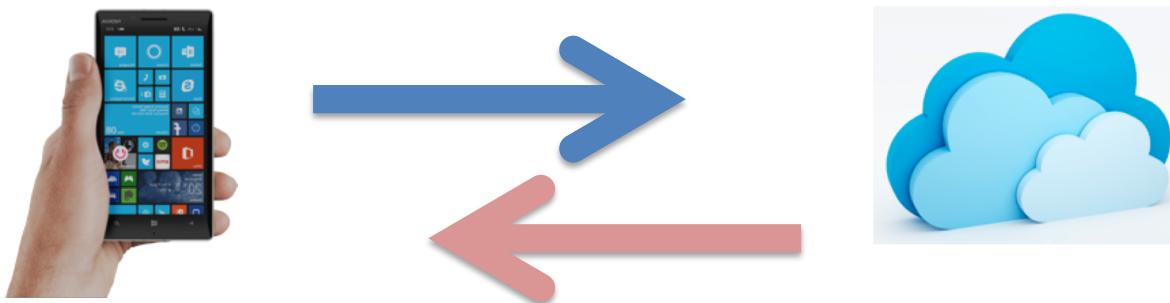


- **Attack on few isolated resources**
 - Soft errors - corruption of part of the computation
- **Massive attacks**

Countermeasures against such attacks (detect/correct)

Verifiable [outsourced] computation

- Trusted but slow Client (Verifier, *Victor*) sends a function F with input x to the server



- Fast but untrusted Server (Prover, *Peggy*) returns $y = F(x)$ and a proof Π that y is correct.

Computing Π should take almost same time than F

Verifying Π should take less time than computing F .

The power of interaction



Verifier
(Victor)



Prover
(Peggy)

Example : Graph isomorphism

- Input : two graphs G_1 and G_2
- Output : Yes iff G_1 and G_2 are isomorphic
- **Graph isomorphism is in NP**
 - not known to be neither in P nor NP-complete;
 - not known to be in co-NP.
- **Proof of isomorphism :** Π = permutation that maps E_1 to E_2
 - Then a verifier can check the proof in almost linear time.
- *But no polynomial proof yet known for graph non-isomorphism.*

Interactive graph [non]-isomorphism

- Victor
 - Toss $b := \text{rand}\{1,2\}$
 - $H := \text{random_permutation}(E_b)$
 - Asks Peggy: to which H is isomorphic to : E_1 or E_2 ?
-
- The diagram illustrates the interaction between Victor and Peggy. On the left, a hand holds a smartphone displaying a Windows interface. A large blue arrow points from the phone to a white cloud icon on the right. Below the blue arrow, the text "Peggy returns y and Π " is written. A red arrow points from the cloud back towards the phone.
- Victor checks Π and if OK
 - If $y \neq b$: Victor has a proof that E_1 isomorphic to E_2
 - Else $y = b$: Victor states that E_1 is not isomorphic to E_2 with error probability $\frac{1}{2}$
 - Error probability $\leq 2^{-k}$ after k rounds

Efficient verifiable computing by spot checking

- Check polynomial equality by random evaluation
 - Choose r_1, \dots, r_n at random in a subset S of a field
 - If $Q(r_1, \dots, r_n) = 0$ then $Q = 0$ with error probability $\leq \deg(Q) / \#S$
- Example: *Verifying matrix multiplication* (**Friedval's algorithm**)
 - To check $C = A \cdot B$, choose a random vector r and verify $C \cdot r = A \cdot (B \cdot r)$

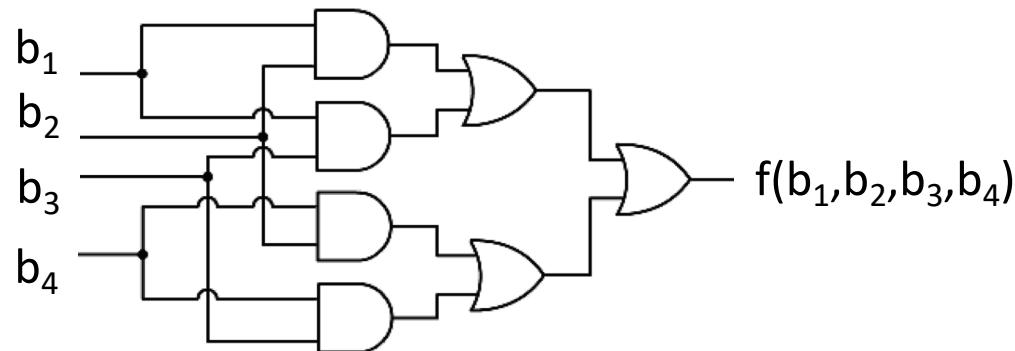
Cost : linear in $\text{size}(A) + \text{size}(B) + \text{size}(C)$

Interactive linear algebra

- Most dense linear algebra reduces to Matrix multiplication
 - Locally compute the (recursive) scheme in $O(n^2)$ while outsourcing all Matrix Multiplications
 - [Algorithm-Based Secure and Fault Tolerant Outsourcing of Matrix Computations, A Kumar, JL Roch, HAL 2013]
- Alternatively provide efficient certificates for sparse linear algebra
 - [Interactive certificates for linear algebra, JD Dumas, E Kaltofen, ISSAC 2014]

Verifying general circuits

- Inputs : $b_1 \dots b_n$ Outputs : $y_1 \dots y_m$
- How to verify $y_1 \dots y_m = f(b_1 \dots b_n)$

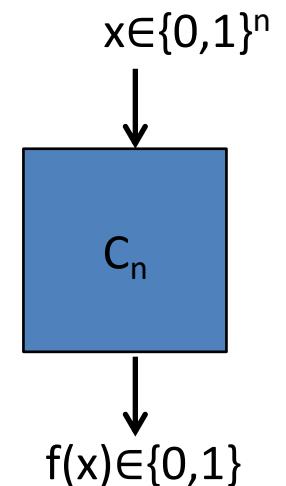
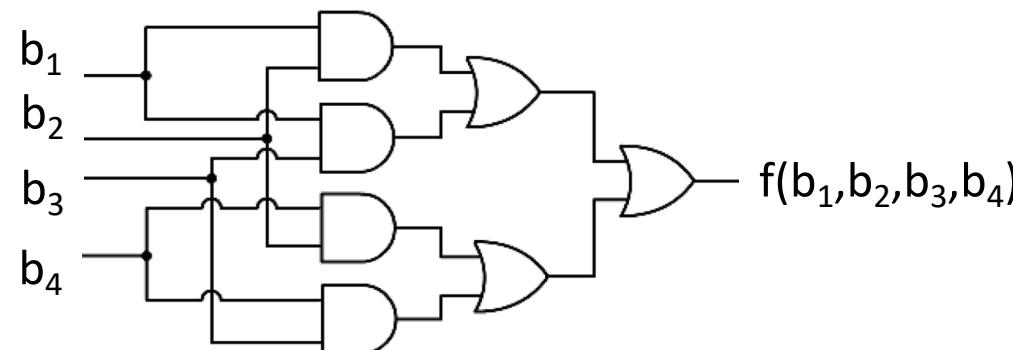


The power of interaction

- Any problem in PSPACE has a polynomial verifier
 - TQBF (quantified Boolean formula problem)
- A polynomial interactive scheme for #SAT

A key tool: the sum-check protocol

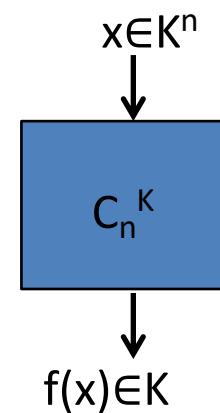
- **Input** : a (boolean) circuit C_n of depth δ that implements a function f with n bits in input :



- **Output** : $S_n = \sum_{b_1=0,1} \dots \sum_{b_n=0,1} f(b_1, \dots, b_n)$
- Let $d=2^\delta$: #usefull gates $\leq d$.
 Theorem: The verifier *interactively* computes S_n in polynomial time $(n+d)^{O(1)}$. (if $\delta=O(\log n)$, polynomial in n)
- Application: number of elements that verify a predicate (#SAT)

Key 1: Arithmetization

- Transform the boolean circuit C_n in an arithmetic circuit C_n^2 in any field K (eg mod p) :
 - $x \text{ and } y = x \cdot_K y$ $\text{not}(x) = 1 - x$
 - $x \text{ or } y = \text{not}(\text{not}(x) \text{ and } \text{not}(y)) = 1 -_K (1 -_K x) \cdot_K (1 -_K y)$
- Transform the circuit C_n^2 in a circuit C_n^K with input in a (large) field K .
 - Gates are $+$ and x in K
 - When inputs are 0 or 1, the output is the same than C_n
- Now, the circuit can be seen as a polynomial in n variables (the input) with degree d
 - For $m=\log \#K$, the circuit can be evaluated in time $(nm)^{O(1)}$, polynomial for any [random] input in K^n .
- **Key 2: induction on the number of sum**
 - Each sub-sum is verified with Schwartz-Zippel



Interactive verification of #3-SAT

- Let: $\Phi = (c_1 \text{ and } \dots \text{ and } c_m)$ be a 3-SAT CNF formula
- Arithmetization of Φ gives $g(X_1, \dots, X_n) = Q(c_1) \cdot Q(c_2) \dots Q(c_m)$
- $\text{Deg}(g) \leq 3m$ (small)
Polynomial-size circuit to evaluate g at any (b_1, \dots, b_n)
- To prove $\#\text{SAT}(\Phi)=K$ reduces to a sequence of sum-check
 - computation in F_p with p prime $> 2^n$

[Goldwasser, Kalai, Rothblum 2008][Thaler Crypto 2015]

Outsourcing general circuits

- Circuits C with n inputs and outputs,
 - Work W , depth D
 - Each level is of degree 1 (multilinear extension)
- Computation is valid iff all levels are corrects
 - Verified by a sum-check at each level
- Cost = $(N + D) \log^{O(1)} (N + W)$
- Optimization when the computation resumes to a reduction of independent parallel computations

Illustration on Matrix Multiplication

[Thaler, Crypto 2013, Crypto 2015]

- Let A and B matrices (n,n) in K with $m = \log_2 n$
- A is a (boolean) function $\{0,1\}^m \times \{0,1\}^m \rightarrow K$:
$$A(i_1, \dots, i_m, j_1, \dots, j_m) = A(i, j)$$
- Let g_A be the polynomial multilinear extension of A
- The g_C verifies
$$g_C(i_1, \dots, i_m, j_1, \dots, j_m) = \sum_{k=0..n} g_A(i_1, \dots, i_m, k_1, \dots, k_m) \cdot g_B(k_1, \dots, k_m, j_1, \dots, j_m)$$
- With the sum-check protocol, this sum of n elements is verified in $O(\log n)$
- Generalizes to parallel computations with logarithmic depth (NC1)

Practical efficiency ?

- Further improvements [Thaler]
 - Sum of products only
 - Same circuit for any coefficient

Problem Size	Naïve MatMult Time	Additional P time	V Time	Rounds	Protocol Comm
1024 x 1024	2.17 s	0.03 s	0.67 s	11	264 bytes
2048 x 2048	18.23 s	0.13 s	2.89 s	12	288 bytes

- Yet far from Fiedvald's verification

Conclusion on outsourcing

- Verifying delegated computation
 - Interaction between models provides power
 - Enables the provable use of untrusted platforms
 - Overclocked processors, algorithms with faults, quantum computing, ...
 - Fully Homomorphic encryption (yet inefficient)
 - Current research to improve FHE efficiency
- On going research - Applications
 - Cloud computing. (web services)
 - Outsourced fault-tolerant computation
 - Secure remote storage (privacy)
 - Secure control-command for critical infrastructure (SCADA)
 - A promising market (eg digital doctor)

Questions ?

- S Arora, B Barak. Computational Complexity: A Modern Approach. Cambridge University Press, 2009
- C. Gentry, Fully homomorphic encryption using ideal lattices, 2009
- J Thaler. Time-optimal interactive proofs for circuit evaluation. Crypto 2013,
- D Fiore, R Gennaro, Publicly Verifiable Delegation of Large Polynomials and Matrix Computations, with Applications, 2013
- B Parno, C Gentry. Pinocchio: nearly practical verifiable computation. Security&Privacy 2013
- M Walish, AJ Blumberg Verifying computations without reexecuting them. Comm. ACM 2/2015
- I Chillotti, N Gama, M Georgieva, M Izabachene. "Improving TFHE: faster packed homomorphic operations and efficient circuit bootstrapping". 2017.